

Fuzzy-Genetic Hybrid Models for Large-Horizon Deterministic Planning

Final Master's Project

Mark Safrónov

10th September 2025

Princess Maker 2 — where it all started



Princess Maker 2 — scale

22 numeric attributes



Princess Maker 2 — scale

25 actions

each changes 2+ attributes



Princess Maker 2 — scale

25 actions

each changes 2+ attributes

Decorum	21	
Art	20	
Conversation	8	
Cooking	19	+1
Cleaning	16	+1
Temperament	23	+1
Charisma	24	
Morality	32	
Faith	20	
Sin	0	
Sensitivity	19	-2
Stress	0	



Princess Maker 2 — endings

61 ending each having its own requirements to meet

Example: Queen ending

Stats Required

Refinement/Elegance: 800+

One of these stats should be 421+:

Fighter Reputation

Magical Reputation


Social Reputation

Housework Reputation


Art must NOT be her highest acquired skill stat

The gap between highest and lowest reputation should be below 50.


She must NOT be eligible for a Sinful Ending (the requirements for the Sinful endings are Morality below 30 and Sin 250+ or alternatively, Morality 0 and Sin 100+)



Blood Type A, Birthday 1200-10-1.	
Libra, VENUS	
Height 154.11cm	Weight 39.68kg
B 85.06cm	W 55.45cm
H 82.75cm	
Physical Fitness.....	808
Strength.....	0
Intelligence.....	547
Elegance.....	923
Glamour.....	333
Morality.....	524
Faith.....	350
Sin.....	0



"Mighty one, I have been watching you. You suit the role of father fairly well. Your daughter has gained some extremely rare abilities. Her education went well!"



Can we solve that?

From the game to a CS problem

Problem: given an ending we want, can we deduce a list of actions to reach this ending?

From the game to a CS problem

One-page definition of the problem

Deterministic planning problem with the fixed-length actions sequence

2.1 Actor behavior as a control problem

Assuming we have a character described as a set of numeric characteristics

$$\mathbf{x} \in \mathbb{Z}^n \quad (2.1)$$

we have a set of possible actions

$$A = \{a_1, a_2, \dots, a_m\} \quad (2.2)$$

which collectively form a transfer function

$$f(\mathbf{x}, a) = \mathbf{x}' \quad (2.3)$$

To describe the desired outcome, we first declare a fitness function mapping the state to a numerical value:

$$\Phi : \mathbf{x}' \rightarrow \mathbb{R} \quad (2.4)$$

a goal fitness value

$$\mathbf{G} \in \mathbb{R} \quad (2.5)$$

and a planning horizon

$$T \in \mathbb{Z} \quad (2.6)$$

We want to get an ordered actions sequence of length T which will lead \mathbf{x} to some \mathbf{x}^* :

$$\mathbf{a} \in A^T, x_o = \mathbf{x} : \bigodot_{i=1}^T f(x, a_i) = \mathbf{x}^* \quad (2.7)$$


(where \odot is a fold operator)

such as:

$$\Phi(\mathbf{x}^*) > \mathbf{G} \quad (2.8)$$

The problem of scale

Depth of the tree is **360+** steps, depending on how do we interpret scheduling



Every fan-out is 25 options not 2

Bruteforce enumeration is intractable, what are our other options?

Option 1: Automatic planning

Courtesy of Dr. Hatem Abdelatiff.

There exist ready-made solutions based on STRIPS formalism in the PDDL syntax.

Solver like ENHSP can work with numeric states and goals.

It will not work for us because **we must have fixed length plans** — defining a cost function to enforce this is nontrivial if possible at all.

Option 2: Reinforcement learning

Courtesy of Dr. Jordi Duch.

«Stateful actor learns to perform sequences of state-changing actions to reach predefined goals»

It will not work for us because of a **scale and delayed evaluation** — after 360...1200th step the signal will be negligible.

Treat the problem as a control problem

What if I treat every trajectory as a control problem.

Then I need only to invent a way to influence the decisions in that single trajectory, and optimise on this «way»

2.1 Actor behavior as a control problem

Assuming we have a character described as a set of numeric characteristics

$$\mathbf{x} \in \mathbb{Z}^n \quad (2.1)$$

we have a set of possible actions

$$A = \{a_1, a_2, \dots, a_m\} \quad (2.2)$$

which collectively form a transfer function

$$f(\mathbf{x}, a) = \mathbf{x}' \quad (2.3)$$

To describe the desired outcome, we first declare a fitness function mapping the state to a numerical value:

$$\Phi : \mathbf{x}' \rightarrow \mathbb{R} \quad (2.4)$$

a goal fitness value

$$\mathbf{G} \in \mathbb{R} \quad (2.5)$$

and a planning horizon

$$T \in \mathbb{Z} \quad (2.6)$$

We want to get an ordered actions sequence of length T which will lead \mathbf{x} to some \mathbf{x}^* :

$$\mathbf{a} \in A^T, x_o = \mathbf{x} : \bigodot_{i=1}^T f(x, a_i) = \mathbf{x}^* \quad (2.7)$$

(where \odot is a fold operator)

such as:

$$\Phi(\mathbf{x}^*) > \mathbf{G} \quad (2.8)$$

Clustering using domain knowledge

«Inclinations»

Fighting	Magic	Artistry	Housekeeping	Sinfulness
<div>Constitution Strength</div> <div>Combat Skill Combat Attack Combat Defense</div>	<div>Intelligence</div> <div>Magical Skill Magical Attack Magical Defense</div>	<div>Refinement Charisma</div> <div>Sensitivity</div> <div>Decorum Art Conversation</div>	<div>Morality Faith</div> <div>Cooking Cleaning Temperament</div>	<div>Sin</div>
Lumberjack, Combat classes, etc	Magic classes, Graveyard etc	Dance classes, Tutoring etc	Innkeeping, Theology etc	Cabaret etc

Single trajectory

5 inclination values + current state at step n controls the priorities of actions



Selected highest-priority action changes the current state



Repeating T times gives us a trajectory and a final state



Evaluate the final state (whether we are at a goal)

Optimization

5 inclination values effectively define the final fitness



We have a pure function to *argmin*

Chosen approaches

Inclinations + current state must map to action priorities: **fuzzy logic**

We must optimize inclinations to minimize fitness: **evolutionary computations**

(minimize not maximize fitness is a technical detail)

Fuzzy logic

Strength 215

→ «low»: 0.15, «high»: 0.65

Refinement 266

→ «low»: 0.12, «high»: 0.67

```
rule: if InclinationFighting is high and strength is low then Lumberjack is high
rule: if InclinationArtistry is high and refinement is low then Lumberjack is low
```

Lumberjack priority: **0.11**

(just an example, depends
on the defuzzification)



Lumberjack priority «high»: 0.1, «low»: 0.1

Evolutionary computations

5 inclinations: vector (I1, I2, I3, I4, I5) — «**genome**»

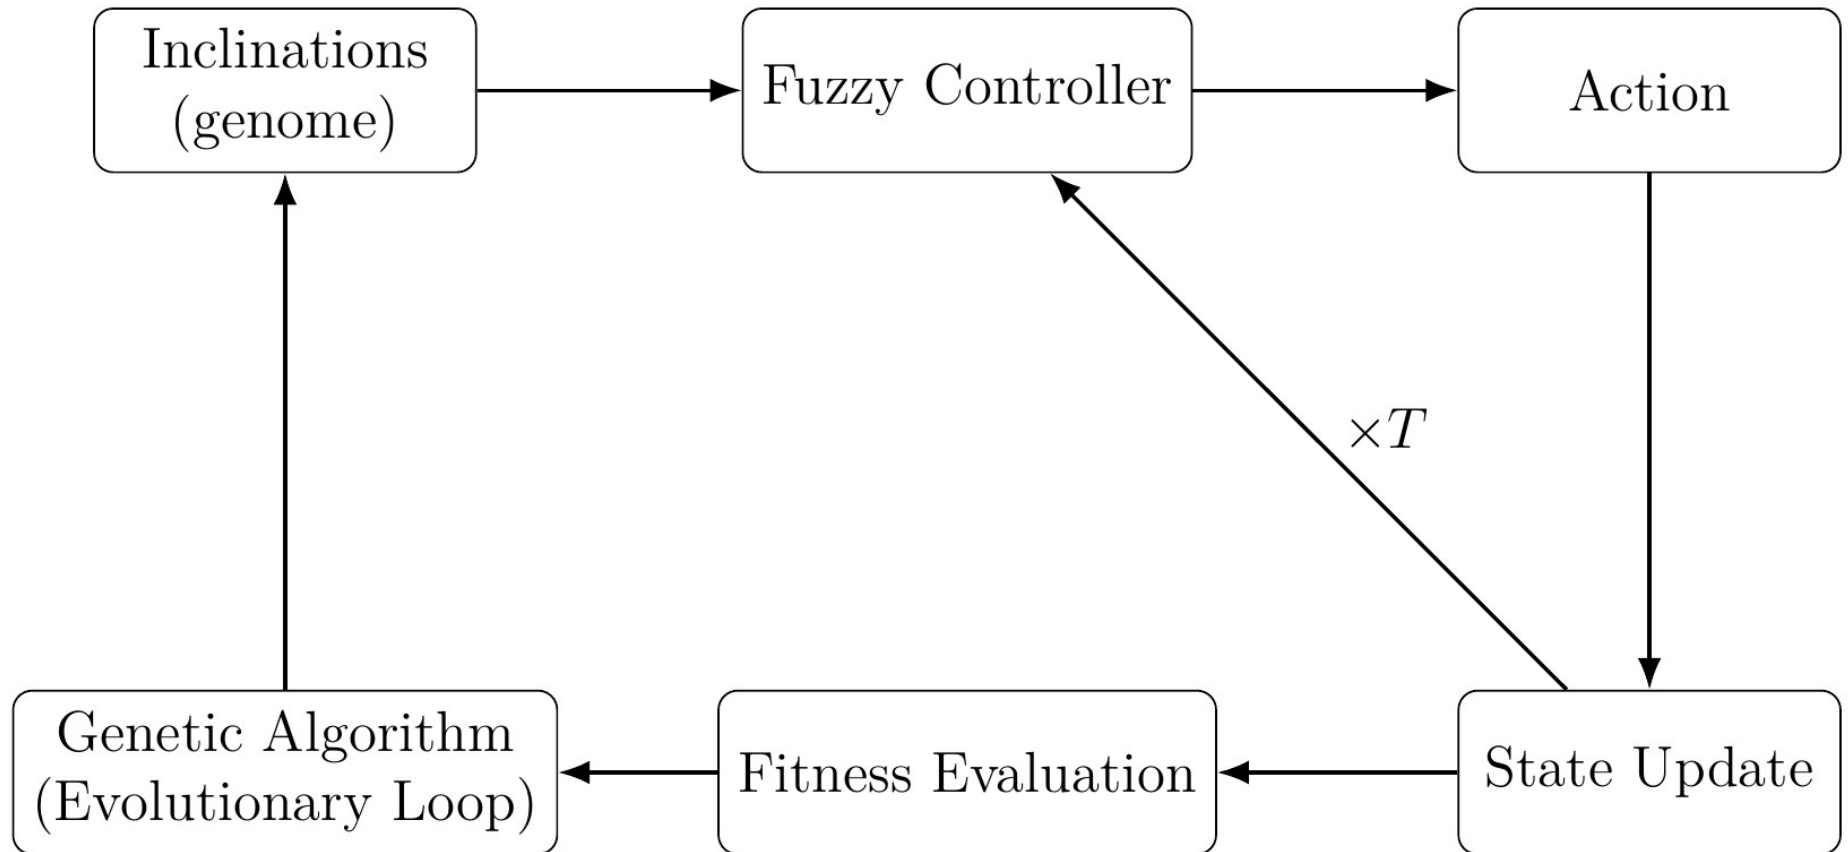
Generate N random genomes — «**population**»

For every genome, calculate their fitness.

Apply special evolution-inspired «**crossover**», «**mutation**», «**selection**» and «**replacement**» operators to the population

Repeat. If your selection and replacement are sound, computation converges.

Control flow scheme



GOALS

1. Build a proof-of-concept
2. Ensure it actually produces solutions (fitness = 0)
3. Estimate performance
4. Summarize potential problems and possible future research

Baseline experiment setup

Number of character attributes: 21

Number of actions: 25

Length of actions sequence: 1200

Complex goal predicate based on the actual game («High General ending»)

Result:

```
Stats changes:  
str: +0, con: +0, int: +507,  
ref: +0, cha: +0, mor: +102,  
fai: +612, sin: -204, sen: -588  
cs: +330, ca: +165, cd: +165,  
ms: +0, ma: +0, md: +408,  
dec: +0, art: +0, elo: +0,  
coo: +294, cle: +294, tem: +294  
Simulation Result:  
Fitness: 0
```

RESULTS

Linear time on T, depends only on the complexity of the goal predicate

```
1 double fitness(const Stats& stats)
2 {
3     const int fighter_reputation = std::get<0>(stats) + std::get<1>(stats)
4     + std::get<9>(stats) + std::get<10>(stats) + std::get<11>(stats);
5
6     // if sensitivity is higher than intelligence or faith,
7     // return an absolute instant loss
8     if (std::get<8>(stats) >= std::get<2>(stats)
9         || std::get<8>(stats) >= std::get<6>(stats)) {
10         return std::numeric_limits<double>::max();
11     }
12
13     double fitness_value = 0.0;
14
15     // add penalty for intelligence below 500
16     if (std::get<2>(stats) < 500) {
17         fitness_value += (500 - std::get<2>(stats));
18     }
19
20     // add penalty for morality below 30
21     if (std::get<5>(stats) < 30) {
22         fitness_value += (30 - std::get<5>(stats));
23     }
24
25     // add penalty for faith below 300
26     if (std::get<6>(stats) < 300) {
27         fitness_value += (300 - std::get<6>(stats));
28     }
29
30     // add penalty for fighter reputation below 421
31     if (fighter_reputation < 421) {
32         fitness_value += (421 - fighter_reputation);
33     }
34
35     return fitness_value;
36 }
```

Stats changes:

str: +0, con: +0, int: +507,
ref: +0, cha: +0, mor: +102,
fai: +612, sin: -204, sen: -588
cs: +330, ca: +165, cd: +165,
ms: +0, ma: +0, md: +408,
dec: +0, art: +0, elo: +0,
coo: +294, cle: +294, tem: +294

Simulation Result:

Fitness: 0

Requirements clauses amount	Runtime on a reference machine
full "General" ending goal	12 m 30 s
no reputation check	6 m 30 s
no reputation, no faith checks	5 m 50 s
no reputation, no faith, no morality checks	6 m 10 s
no attribute checks, only the predicate	3 m 20 s

RESULTS

Result coming out of the solver is a vector of inclinations (5 real numbers in our setup).

Because the problem is formulated completely deterministic, we can effectively decode the sequence of actions out of these inclination values.

As such, our proof-of-concept deduces action sequences of fitness 0 \rightarrow our goals are reached.

1. TheologyClass x 383
2. FencingClass x 106
3. FightingClass x 106
4. TheologyClass - FencingClass - FightingClass x 24
5. TheologyClass
6. FencingClass - FightingClass x 2
7. TheologyClass - FencingClass - FightingClass x 33
8. Housework x 294
9. Church x 63
10. Church - ScienceClass x 3
11. Church
12. Church - ScienceClass x 6
13. Church
14. Church - ScienceClass x 7
15. Church
16. Church - ScienceClass x 6
17. Church
18. Church - ScienceClass x 7
19. Church
20. Church - ScienceClass x 4
21. Church

Conclusions

- 1) A generic problem has been formulated and a proof-of-concept solver has been built for it.
- 2) We successfully utilized domain knowledge to reduce dimensionality of the original problem.
- 3) On the given instance of a problem the solver reaches linear time over length of action sequence T . Complexity of the goal predicate significantly influences the runtime.
- 4) Still yet to research the dependence on number of actions and number of attributes.
- 5) Still yet to compare with the reference possible implementations of solvers using ENHSP and reinforcement learning.

Thank you for listening