

UNIVERSITAT ROVIRA I VIRGILI

MASTER'S THESIS

MASTERS OF ARTIFICIAL INTELLIGENCE AND COMPUTER SECURITY

Fuzzy-Genetic Hybrid Models for Large-Horizon Deterministic Planning

MARK SAFRONOV

September 2, 2025



UNIVERSITAT
ROVIRA I VIRGILI

Contents

1	Introduction	3
1.1	Motivation and Context	3
1.2	Existing Approaches	4
1.2.1	Automatic Planning Theory	4
1.2.2	Reinforcement learning	4
1.3	Fuzzy Logic and Fuzzy Controllers	5
1.4	Genetic Algorithms	5
1.5	Our Proposed Approach	6
1.6	Comparative assessment	7
1.7	Contributions and Objectives	8
2	Formal problem statement	10
2.1	Actor behavior as a control problem	10
2.2	Dimensionality explosion stemming from the original context	11
3	Proposed solution approach	12
3.1	Using the domain knowledge to reduce dimensionality	12
3.2	Hypothesis	13
3.3	Objectives	13
4	Methodology	15
4.1	Encoding domain knowledge of actions as a fuzzy controller	15
4.2	Control feedback loop as a fitness function	17
4.3	Global optimization using an evolutionary algorithm	18
5	Implementation	19
5.1	Choice of a C++ language as foundation	19
5.2	Fuzzylite library for the fuzzy controller implementation	19
5.3	Pagmo library for evolutionary computations	21
5.4	Implementing the method using fuzzylite and pagmo libraries	22
5.5	Using the implemented solver	26
6	Experiment 1: Trivial case	28
6.1	Discussion of the results	32
7	Experiment 2: Base control case	34
7.1	Discussion of the results	38

8	Conclusions and Future Work	42
A	Baseline case fuzzy controller	47

Chapter 1

Introduction

1.1 Motivation and Context

A deterministic planning problem for a stateful agent is one of the most generic formulations in artificial intelligence. The task can be summarized as selecting a sequence of actions which transforms the state of an agent into a desired goal state. This formulation naturally appears in a wide range of real-world domains: robotics (planning robot trajectories under resource and time constraints), industrial scheduling (allocating tasks to maximize production while minimizing downtime), and control systems (maintaining stability of dynamic processes under changing conditions).

The key difficulty in such problems is the combinatorial explosion. As soon as the number of available actions grows and the planning horizon becomes long, the search space of all possible trajectories expands exponentially. This phenomenon, commonly referred to as the curse of dimensionality, renders classical exhaustive approaches infeasible even for moderately sized problems.

Works like [4] cover both the real-world applications and a significance of the curse of dimensionality.

In this work, we use the Japanese computer game *Princess Maker 2* (Gainax, 1993) as a motivating case study. While the game itself is narrative-driven and contains many mechanics that go beyond formal modeling, its core gameplay loop is a textbook instance of a deterministic planning problem. The player controls a stateful agent — the daughter character — by repeatedly selecting actions such as study, work, or leisure. Each action deterministically modifies a set of numerical attributes, and after a fixed horizon of in-game years, the outcome is evaluated against predefined conditions that determine the “ending.”

Stripped of its narrative elements, *Princess Maker 2* yields a highly formal planning problem: an agent with dozens of numeric attributes, several dozen possible actions, and a planning horizon of several thousand steps. This combination of a large state space, broad action set, and long horizon exemplifies the type of problem where classical planners and reinforcement learning approaches both struggle. At the same time, because the formulation is generic, any progress on this problem carries implications well beyond the original game context.

In the following section we will briefly review the most obvious approaches to this problem, namely, automatic planning theory and reinforcement learning, and discuss

why they both fail to scale to the case we are interested in.

1.2 Existing Approaches

1.2.1 Automatic Planning Theory

Automatic planning in artificial intelligence is historically grounded in the STRIPS [6] formalism, which defines planning as the composition of operators that transform an initial world model into one satisfying a goal condition. In principle, this description matches exactly the type of solver we are aiming for: we have a stateful agent, a set of deterministic actions, and a requirement to reach a specified goal state.

Modern planners, such as ENHSP [14] [13], extend this formalism to support numeric fluents and continuous effects. On small-scale problems they have been demonstrated to work effectively, often guaranteeing optimality of the resulting plan. For example, existing benchmarks ¹ include instances with up to forty numeric variables and around ten actions, which can be solved in practical time.

However, the core assumptions of classical planning diverge from the requirements of our case. Most planners are designed to minimize plan length, i.e., to find the shortest action sequence that satisfies the goal. In contrast, the Princess Maker 2 formulation requires executing a fixed number of steps before the final state can be evaluated. No-operation actions are not available, so every step transforms the agent’s state irreversibly. This means that reaching the goal “too early” is not valid — it is only the state after the full horizon that determines success.

In addition, the problem of scale remains prohibitive. Even if the goal condition could be reformulated to suit the planner’s semantics, the state space induced by 23 attributes, 27 actions, and thousands of steps is far outside the scope of tractable automatic planning. Research does exist on scaling planners with temporal logics [8] and heuristics, but these approaches are not focused on high-dimensional numeric fluents of the type present in our motivating case.

1.2.2 Reinforcement learning

Reinforcement learning [18] (RL) is another approach which seems, at first glance, to be a natural fit for the problem. In RL, we model an agent interacting with an environment: the agent observes a state, selects an action, and receives a reward. Over time, the agent is expected to learn a policy that maximizes cumulative reward. The analogy with our setting is straightforward — the state is the character’s attributes, the actions are the available schedule choices, and the “reward” is the final outcome at the end of the planning horizon.

However, the practical obstacles are significant. The first is the dimensionality of the state space. With 23 attributes each taking values in the range of roughly 0 to 500, the number of possible states is astronomically large. A naive tabular representation of a Q-function would require a table of size 27×500^{23} , which is intractable to store or

¹<https://github.com/hstairs/enhsp/tree/enhsp-20>

update. While function approximation can reduce this requirement, it does not remove the exponential blowup inherent in such spaces.

The second issue is the length of the planning horizon. In the simplified gameplay setup, the player makes three choices per month for ten in-game years, which already leads to a search tree over 360 steps. In the full day-by-day formulation the horizon is an order of magnitude longer. Standard RL algorithms struggle when the reward is sparse and delayed, as in this case where only the terminal state matters. Exploration becomes ineffective because the probability of randomly reaching a good ending within thousands of steps is vanishingly small.

Finally, unlike in many RL benchmarks, our problem does not allow neutral or reversible actions. Every action irreversibly changes the state, which exacerbates the exploration problem. Taken together, these aspects make classical reinforcement learning algorithms impractical for the problem scale we consider.

1.3 Fuzzy Logic and Fuzzy Controllers

Fuzzy logic, introduced by L. Zadeh in 1965, extends classical set theory by allowing partial membership of elements in a set. Instead of a variable belonging to a set either completely or not at all, in fuzzy logic a variable may belong to a set with a degree between 0 and 1. This allows us to model linguistic terms such as “low,” “medium,” or “high” in a mathematically precise way.

A fuzzy controller applies this idea to decision-making. The system is defined through linguistic variables with associated membership functions, and a collection of symbolic rules. A typical example outside this work would be:

`If temperature is high then fan speed is high.`

Here, both “temperature” and “fan speed” are fuzzy variables, and the rule connects them in a way that corresponds directly to human reasoning. After applying fuzzification (mapping numeric input values into fuzzy terms), rule evaluation, and defuzzification (converting fuzzy outputs back into numeric values), the controller produces an actionable decision.

The benefit of this approach is the ability to capture domain knowledge without enumerating every possible combination of numeric parameters. Instead of constructing a complete reward table or transition map — which would be infeasible in large state spaces — we specify a relatively small set of symbolic rules which generalize across many states. This makes fuzzy controllers attractive for problems where expert intuition about the domain exists, but exhaustive specification is impossible.

In the context of this work, fuzzy controllers provide a mechanism to express priorities over actions in symbolic terms, forming a bridge between human-understandable heuristics and algorithmic decision-making.

1.4 Genetic Algorithms

Genetic algorithms [16] [1] (GAs) are a class of optimization methods inspired by biological evolution. A population of candidate solutions is maintained, where each solution

is represented as a genome — a sequence of values encoding the parameters of interest. The algorithm proceeds iteratively through the following steps:

1. Evaluation: every genome is assigned a fitness score according to an objective function.
2. Selection: genomes with higher fitness are more likely to be chosen for reproduction.
3. Crossover: pairs of genomes exchange parts of their structure, creating new offspring that combine traits of both parents.
4. Mutation: with low probability, individual genes are randomly altered, introducing diversity.
5. Replacement: the new generation replaces the old, and the process repeats.

Over time, the population tends to evolve towards fitter solutions, even when the search space is large, non-linear, or poorly understood. The method does not guarantee a globally optimal solution, but it often finds good approximations in domains where exact methods are computationally infeasible.

The main advantage of GAs is that they treat the problem as a black box: only the ability to evaluate candidate solutions is required, while no assumptions about continuity, differentiability, or convexity are needed. This makes them well suited for combinatorial or highly non-convex problems.

In the context of this work, genetic algorithms will be used to optimize the inclinations — latent behavioral parameters that guide the fuzzy controller. This reduces the original planning problem from direct search over long sequences of actions to an evolutionary search in a much smaller parameter space.

1.5 Our Proposed Approach

The discussion above shows that classical approaches such as automatic planning or reinforcement learning are unable to scale to the requirements of our motivating problem. To address this, we propose a hybrid heuristic method which combines symbolic reasoning with evolutionary optimization.

The core idea is to replace direct search over complete action sequences with search over a smaller set of latent behavioral parameters, which we call inclinations. These inclinations represent tendencies toward particular groups of actions. A fuzzy controller then maps the inclinations together with the current state into a concrete choice of action.

This formulation has two benefits. First, it allows us to encode domain knowledge directly into the fuzzy rule base, reducing the effective dimensionality of the search. Second, it enables optimization to take place in the smaller inclination space, rather than in the exponentially large space of all possible trajectories. The resulting search problem becomes tractable for evolutionary algorithms such as genetic algorithms, which only require evaluation of candidate solutions, not explicit enumeration of the search space.

The hypothesis of this work is that such a hybrid fuzzy-genetic approach can find valid solutions in polynomial-like time in cases where reinforcement learning or automatic planners fail due to combinatorial explosion.

The overall process is illustrated in Figure 1.1.

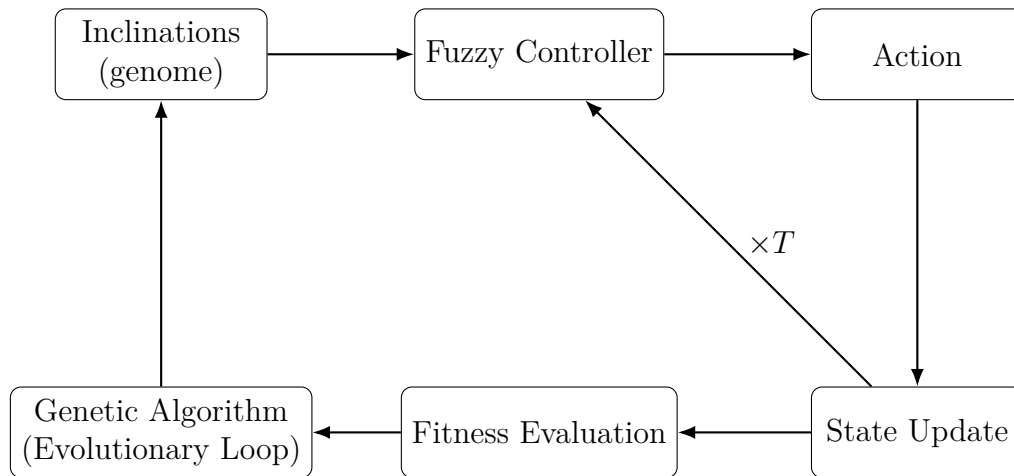


Figure 1.1: Schematic overview of the fuzzy-genetic approach.

1.6 Comparative assessment

To summarize the discussion above, in the Table 1.6 we compare the two obvious alternatives with the hybrid fuzzy-genetic approach proposed in this work.

Method	Strengths	Weaknesses in our context
Automatic planning theory	<ul style="list-style-type: none"> - Formal match to state-action-goal formulation. - Established solvers (e.g., ENHSP) exist. - Optimality guarantees on small problems. 	<ul style="list-style-type: none"> - Focus on minimizing plan length, incompatible with fixed-horizon requirement. - No-operation actions not available, premature goal achievement invalidates plan. - State space scaling beyond practical limits.
Reinforcement learning	<ul style="list-style-type: none"> - Natural agent/action/reward formalism. - Flexible and general framework. - Large body of research and practical algorithms. 	<ul style="list-style-type: none"> - Curse of dimensionality: 500^{23} states in the motivating case. - Very long horizons with sparse terminal rewards. - No reversible or neutral actions, exploration ineffective.
Fuzzy-genetic hybrid approach (proposed)	<ul style="list-style-type: none"> - Encodes domain knowledge symbolically via fuzzy rules. - Dimensionality reduction through latent variables (“inclinations”). - Optimization operates on reduced space, not on raw action sequences. 	<ul style="list-style-type: none"> - Manual design of fuzzy rule base is non-trivial. - No global optimality guarantees. - Performance depends on quality of rules and evolutionary hyperparameters.

Table 1.1: Comparison of alternative methods for solving large-horizon deterministic planning.

This comparison highlights why classical planners and reinforcement learning approaches are unsuitable for the case we focus on, and motivates the introduction of a fuzzy-genetic heuristic, which will be presented in detail in Chapter 3.

1.7 Contributions and Objectives

1. Formalization of Princess Maker-style problems as deterministic planning with large horizons.
2. Introduction of inclinations + fuzzy rule base as a dimensionality reduction heuristic.
3. Implementation of a solver using fuzzylite (fuzzy logic) and pagmo (evolutionary algorithms).
4. Experimental evaluation on trivial and baseline control cases.

This work is foundational proof of concept, not yet a full game solver nor a generic ready-to-use solver for a wide range of real-world problems.

In Chapter 2 we now give a precise problem formulation, and in Chapter 5 we develop the proposed fuzzy-genetic heuristic approach in detail.

Chapter 2

Formal problem statement

As follows from the section ??, formally we have a choice of whether to treat this as a planning problem or a control problem. Despite Reinforcement learning not being fit for our cause, we will still approach our origin problem as a control problem, as it allows us to formally introduce our solution approach later in the chapter 3.

2.1 Actor behavior as a control problem

Assuming we have a character described as a set of numeric characteristics

$$\mathbf{x} \in \mathbb{Z}^n \quad (2.1)$$

we have a set of possible actions

$$A = \{a_1, a_2, \dots, a_m\} \quad (2.2)$$

which collectively form a transfer function

$$f(\mathbf{x}, a) = \mathbf{x}' \quad (2.3)$$

To describe the desired outcome, we first declare a fitness function mapping the state to a numerical value:

$$\Phi : \mathbf{x}' \rightarrow \mathbb{R} \quad (2.4)$$

a goal fitness value

$$\mathbf{G} \in \mathbb{R} \quad (2.5)$$

and a planning horizon

$$T \in \mathbb{Z} \quad (2.6)$$

We want to get an ordered actions sequence of length T which will lead \mathbf{x} to some \mathbf{x}^* :

$$\mathbf{a} \in A^T, x_o = \mathbf{x} : \bigodot_{i=1}^T f(x, a_i) = \mathbf{x}^* \quad (2.7)$$

(where \odot is a fold operator)
such as:

$$\Phi(\mathbf{x}^*) > \mathbf{G} \quad (2.8)$$

The transfer function f is assumed to be completely determined, and the whole process being non-stochastic. This is a significant restriction which cannot be lifted for the proposed solution to work.

2.2 Dimensionality explosion stemming from the original context

We assume a fixed-length trajectory of actions, each of which transforms the state of the system according to a known deterministic transfer function (2.3).

This means two restrictions:

1. No-operation actions are prohibited, each step must result in a meaningful state transformation, reflecting the irreversible nature of time.
2. Goal state must still be in effect at the step T .

While the agent cannot avoid taking actions — and hence cannot avoid changes to the system — it is allowed to evaluate its progress toward the goal at every intermediate state. In this sense, the problem is not a pure planning task but an episode-based control problem with delayed evaluation.

In this work we'll focus specifically on the cases which lead to combinatorial explosion for classical solutions, that is, when we have sufficiently large amount of characteristics, actions to choose from and most importantly, very large planning horizon:

$$n > 25 \quad (2.9)$$

$$m > 24 \quad (2.10)$$

$$T > 360 \quad (2.11)$$

The origin *Princess Maker* problem is the lower edge of the cases we are interested in.

With these restrictions in place, a need in an heuristic arises to perform efficient search in the state space, as its size becomes unrealistically large.

Chapter 3

Proposed solution approach

Classical reinforcement learning methods become intractable in this domain due to the high dimensionality of the state space, large action set, and long planning horizon. Moreover, the inability to halt or take neutral actions further exacerbates the combinatorial explosion of the trajectory space.

To address this, we introduce a heuristic dimensionality reduction via the concept of inclinations — latent behavioral parameters — and model the behavior policy as a fuzzy controller which maps the current state and inclinations to a concrete action.

This parametrization constrains the space of possible behaviors, making the optimization tractable. Instead of learning or searching over action sequences directly, we perform optimization in the significantly smaller space of inclinations, evaluating the final outcome after T steps. The resulting problem becomes an offline, black-box control task — suitable for evolutionary algorithms, rather than classical RL or automatic planning methods.

3.1 Using the domain knowledge to reduce dimensionality

In this work we evaluate an approach which is defined as follows.

Let's assume that we can segment the set of possible actions to clusters with the following particularities:

1. actions in the same cluster lead to “similar” changes in the character state \mathbf{x} .
2. the cluster as a whole can be described symbolically

In this case we can synthesize a set of numeric characteristics which we'll call “inclinations”:

$$\mathbf{I} \in \mathbb{Z}^q \tag{3.1}$$

$$q \lll n \tag{3.2}$$

From this, we can define a set of fuzzy rules[12] mapping the inclinations to action choices:

1. if an inclination I_i has a fuzzy value V_I ,
2. and the current state \mathbf{x} has fuzzy values V_i^x
3. then P_a , the priority of an action a , is a fuzzy set V_A .

After the defuzzification of all the inferred fuzzy values P_a we select an action with the highest priority.

The selection and design of fuzzy rules is a critical aspect of this approach. In this thesis, the fuzzy rule base is constructed manually, leveraging domain knowledge to define the mapping from inclinations to action priorities. Future research may investigate automated methods for generating fuzzy rules, such as clustering or machine learning techniques, to further improve scalability and reduce manual effort.

While it is theoretically possible to define fuzzy rules that map every possible inclination vector \mathbf{I} or even every state \mathbf{x} to action priorities, such exhaustive rule sets would quickly become infeasible due to combinatorial growth. This reinforces the importance of dimensionality reduction and clustering in making the fuzzy-genetic approach tractable for high-dimensional planning problems.

The assumption which we explore among others in this work is the practical possibility to write a coherent set of fuzzy rules which will be clustered around the clusters of actions, and each inclination will tend to map to its own cluster of actions.

Now, using such a fuzzy controller $\xi(I, \mathbf{x})$ we can construct the goal function:

$$g(I, \mathbf{x}) = \bigodot_{i=1}^T f(x, \xi(I, x_i)) \quad (3.3)$$

the above formula being subject to improvements in expressiveness, the main point of which being the fuzzy controller $\xi(I, x_i)$ selecting the action to perform on the step i according to the inclinations and (ideally) the current state x_i .

The argument \mathbf{x} is essentially a constant for both (2.7) and (3.3). As the transfer function f is non-stochastic, \mathbf{I} uniquely maps to the actions sequence \mathbf{a} . Thus, given (3.2), we effectively performed dimensionality reduction on the original problem.

We can find $\arg \max(g)$ now using an appropriate optimization method. For this work, because of a strong biosocial analogies a genetic algorithm[10] was chosen, with the vector of inclinations \mathbf{I} as a chromosome (every inclination value being a separate gene).

3.2 Hypothesis

The hypothesis explored in this work is that the combination of assumptions described above constructs an heuristic which allows finding a locally optimal solution for the problem defined in chapter 2 in a polynomial time.

3.3 Objectives

This is a foundational work which proves a concept.

That is, whether a fuzzy-genetic heuristic can effectively solve high-dimensional deterministic planning problems through dimensionality reduction and symbolic reasoning.

In particular, we aim to:

1. Formalize the problem as an optimization task.
2. Implement a working solver.
3. Evaluate the performance of the solver on a set of test cases of increasing complexity.
4. Analyze the results to draw conclusions about the effectiveness of the approach.

An objective 1 has been reached in the chapter 2.

As soon as our solver will be able to produce solutions with the fitness representing reaching the goal state, an objective 2 will be considered reached.

An objective 3 is covered by the chapters 6 and 7. Finally, chapter 8 covers the analysis of the results and conclusions about its applicability.

Chapter 4

Methodology

In this section we will discuss the theory which this work is build upon, namely, fuzzy logic [12] and evolutionary algorithms [10].

4.1 Encoding domain knowledge of actions as a fuzzy controller

Normally the Fuzzy logic is being explained from the fuzzy set theory by L. Zadeh [20], but for this particular work the most important part of the fuzzy logic is the fuzzy rules for the fuzzy controller so it's more beneficial to start with them.

In the scope of the FL it is possible to express the domain knowledge in the form of symbolic rules, with the general form as follows:

```
If (input variable A) has a (fuzzy value Fa)
then (output variable B) has a (fuzzy value Fb)
```

For example, for our particular problem and solution method:

```
If InclinationAggressiveness is High
then DuelingClassesPriority is High
```

This rules format depends on the concept of the **Fuzzy Variable**, which is a combination of four major parts:

1. Name
2. Range of “strict” values
3. “strict” value itself
4. Set of fuzzy sets describing the possible fuzzy values of this variable

The concept of Fuzzy Variable, in turn, depends on the concept of a **fuzzy value**, which is a combination of two major parts:

1. Name
2. Membership function

Where the **membership function** is a continuous function mapping the input “strict” values to real numbers between 0 and 1. The *membership function* of a fuzzy value describes the *measure of belonging* of the current “strict” value of the variable to the given symbolic **term**, for example, “high”, “low” and such. Because of the *terms* being literally words from a natural language, fuzzy variable is also called a **linguistic variable**.

Let’s give **an example**. Assume the following fuzzy variable:

$$S = (N, R, V, T) \quad (4.1)$$

$$N = \text{“Strength”} \quad (4.2)$$

$$R = \mathbb{Z} \in [0, 100] \quad (4.3)$$

$$V \in R \quad (4.4)$$

$$T = ((\text{“Low”}, f_l), (\text{“Acceptable”}, f_a), (\text{“High”}, f_h)) \quad (4.5)$$

It specifies three *fuzzy terms* for the numeric property “Strength”, which can have integer “strict” values from 0 to 100. Thus, when we measure this property and provide a strict value for “Strength”, we can determine the values of *membership functions* of its three *fuzzy terms*. For example, if $V = 72$ then

$$T = ((\text{“Low”}, f_l(72)), (\text{“Acceptable”}, f_a(72)), (\text{“High”}, f_h(72))) \quad (4.6)$$

Which should be interpreted as “Strength” of 72 being at the same time $f_l(72)$ “Low”, $f_a(72)$ “Acceptable” and $f_h(72)$ “High”.

The process of calculating the values of membership functions for all the terms of a linguistic variable given its strict value is called **fuzzification** of this value.

The main point of the fuzzification, which we exploit in our method and which is at the core of the fuzzy control theory, is that we get the formal mechanism of transforming numeric values to domain-specific inexact vocabulary.

Fuzzy logic provides the reverse process as well. It is possible to specify the values of the membership functions of all the terms in T of the fuzzy variable, and from them calculate the “strict” value V . This process is called **defuzzification** of the linguistic variable.

Continuing the above example, we can start by specifying the fuzzy values of “Strength” first, possibly, if we measure it by some inexact vague means: $f_l = 0.4$, $f_a = 0.8$, $f_h = 0$.

Then, depending on the exact shape of the functions f_l , f_a and f_h defuzzification gives us a strict value of “Strength”, say, 42.

Given all the above, a **fuzzy controller** is an algorithm which performs three large steps:

1. Applies fuzzification of the values of all the input variables (antecedents of the fuzzy rules)
2. Evaluate all the fuzzy rules, obtaining the fuzzy values of the output fuzzy variables

3. Applies defuzzification to the output fuzzy variables, obtaining their strict values.

The above algorithm is called a **Mamdani fuzzy controller** [9] and it's the one which we'll use in this work.

In our system, we're going to have the vector of inclinations and the current state of the specimen as input variables for the controller, and have the priorities of actions as the output variables. This will allow us to imitate the process of "decision making" of the specimen to choose the next action to perform.

The major benefit and the core reason for the fuzzy controller is the ability to encode the domain knowledge in a limited set of rules which will be formally processed.

Compared to, for example, some of the reinforcement learning methods, we don't need to specify the full table of rewards for every possible action-state combination. It is enough to specify one rule for every available action and the controller will already become fully functional. With some configuration of rules it's possible to write even less of them.

This allows to simplify the implementation of the solver, because one of the main weaknesses of the proposed solution is writing the fuzzy rules by hand.

The second benefit of using the fuzzy controller for decision making is that it can be applied without major changes to non-deterministic, stochastic environment, for example, if the actions would be allowed to make randomized changes to the specimen's state, that is, if the transfer function would not be pure. It opens up the possibilities to explore this topic further in the later works.

4.2 Control feedback loop as a fitness function

A single trajectory in the action space is explored using the following process.

1. We start with the initial state \mathbf{x}_0 and the given set of inclinations \mathbf{I}^k
2. We evaluate both \mathbf{x}_0 and \mathbf{I}^k with the preconfigured fuzzy controller
3. The defuzzified output of the controller is the set of priorities for all the actions. We pick the action with the highest priority. Tiebreaker is the position of the action in the list.
4. Action is executed and if we haven't made T actions yet we return to the step 2
5. After T executed actions we apply the goal conditions predicate $\Phi(\mathbf{x}^*)$ and calculate the fitness based on that.

It is important to understand that the state of a specimen is a transient value, used only for calculations of the final fitness after T iterations. The solution we seek is fully encoded in the inclinations vector \mathbf{I} , which stays unchanged for the entirety of the control loop.

4.3 Global optimization using an evolutionary algorithm

Strong biosocial analogies and the configuration of the control loop from 4.2 suggest us to use the evolutionary algorithms [16] [1] for optimization. This is what would be used in this work. However, in principle, any algorithm which is able to use the concept of a fitness function would be applicable here.

Evolutionary algorithms can be explained with an example of the so-called Simple Genetic Algorithm ¹.

SGA operates on the set of **specimens**, each one being a single option in the search space to explore. A specimen is classically a list of characters, which is called literally a **genome**. The whole set of specimens is called a **population**.

In our case, a specimen would be a list of inclination values.

Every specimen in a population is evaluated using the **fitness function**, producing a fitness value.

Then, a **selection operator** is applied, choosing a subset of the population. For example, our selection operator may be choosing the top 50% of the population by their fitness value.

After the selection, we apply the **crossover operator** to the pairs of selected specimens' genomes. The classical crossover operator picks a single place inside both of the genomes and then swaps the resulting halves between them. For example, a genome 'aaaa000' and a genome '1111bbb' after the crossover at point 5 become 'aaaabbb' and '1111000'.

After the crossover we apply the **mutation operator** to all of the selected genomes. The mutation changes (with some low probability) individual genes in the genomes at random. For example, we can have a mutation operator which has 0.01 probability of flipping a gene in the genome from 'a' to 'b' and *vice versa*. Then, we have 0.002 probability of a specimen with a genome 'aaabb' turning into 'ababb'.

After the crossover and mutation, we finally apply the **replacement** which forms the new population for the next generation and the next round of evolution. For example, we can use a so-called $(\mu + \lambda)$ -evolution strategy [15]: calculate the fitness for all the new genomes and then pick s ones with the best fitness from both the old genomes and new ones, where s is the target population size. The size of the population is being kept constant for the classical genetic algorithms, the role of the replacement operator is specifically to enforce that.

In the approach described in this work, the fitness function is the control loop described in the previous section 4.2. The genome of the specimen is the vector of inclinations. And due to the choice of the specific library for the implementation of the evolutionary algorithms we have a wide selection of them, which means, we can explore different options starting from the Simple Genetic Algorithm and continuing with more complicated options.

The library Pagmo [2] includes a lot of already implemented different evolutionary algorithms apart from the simple genetic algorithm so it enables us easier exploration of possibilities in optimizing the full solver.

¹<https://esa.github.io/pagmo2/docs/cpp/algorithms/sga.html>

Chapter 5

Implementation

The technical implementation of the method is performed in C++ [17] using the libraries FuzzyLite [11] and Pagmo [2]

5.1 Choice of a C++ language as foundation

As the root problem of this work is the problem of scale, it has been decided that we trade comfort of experimentation for pure processing power.

Contemporary C++, starting with the standard version 20, allows for a very high-level code as readable as a natural language. It also has libraries for both the fuzzy logic [11] and evolutionary computations [2] for us to not implement any of them from scratch.

In talking on choice of the language for the implementation we cannot avoid comparisons with Python ¹, assumed leader and language of choice for scientific experiments. As has been stated above, it has been conscious decision to trade the ability to make rapid changes in the code, especially the ability to run convenient machinery like Jupyter notebooks ², for the raw processing power. This is because the C++20 and later is expressive enough to be as readable as Python sans some of the required syntax boilerplate, and in reality the most painful part of choosing C++ is building the program to be cross-platform, as Python programs are, and doing that with the code which uses third-party libraries is a nontrivial implementation problem.

5.2 Fuzzylite library for the fuzzy controller implementation

This work turned out to be more or less an assessment of usefulness of the `fuzzylite` [11] C++ library in addition to the main goal. While being fully open sourced with a non-restrictive license terms, actually adding it to an existing C++ program is a task certainly not feasible for an arbitrary computer scientist not being a seasoned software engineer

¹<https://www.python.org/>

²<https://jupyter.org/>

at the same time. Which is a shame, as it offers a straightforward idiomatic API which allows expressing the algorithms in a readable format.

`fuzzylite` also provides a domain-specific language for specifying the fuzzy controller, which allows us to describe this part of the algorithm in a language more expressive than the raw C++ function calls.

On the following code example is a description of a fuzzy variable in the DSL of `fuzzylite`.

```

1  InputVariable: PhysicalInclination
2      enabled: true
3      range: 0 1.000
4      lock-range: false
5      term: tiny Ramp 0.330 0.000
6      term: low Triangle 0.000 0.330 0.670
7      term: high Triangle 0.330 0.670 1.000
8      term: highest Ramp 0.670 1.000

```

Base syntax of this DSL is essentially YAML ³.

At the first line we specify the name of the variable and whether it will be used as an input or an output for the fuzzy rules. Among the properties of the variable we have the numerical range of strict values for it, supplementary flags `enabled` and `lock-range` not interesting at this moment and several `term` declarations which are the concise descriptions of all the linguistic terms of the variable.

In the example only two membership functions are used: `Ramp` and `Triangle`, but `fuzzylite` has around 20 of them predefined at the time of writing this report.

The following line specifies a single term of a fuzzy variable:

```

1  term: tiny Ramp 0.330 0.000

```

In this line, the word `tiny` is the symbolic name of the term, which represents the vague description of the value directly from the domain knowledge.

The word `Ramp` is a keyword selecting the appropriate membership function from among the ones built-in in the `fuzzylite` library. Figure 5.1 displays the plot of this function.

The notation `0.330 0.000` is an internal trick of the library to indicate the downward slope of the ramp by convention instead of some other method. Writing the x values in ascending order would mean that the ramp is increasing instead of decreasing.

Given the definitions of all the fuzzy variables, the list of rules of the fuzzy controller is specified in almost the natural language ⁴:

```

1  RuleBlock: mamdani
2      enabled: true
3      conjunction: Minimum
4      disjunction: Maximum
5      implication: AlgebraicProduct

```

³<https://yaml.org/>

⁴“then” clauses has been moved to the next lines for the line to fit on the paper, in an actual code the rule is written on a single line without breaks




Figure 5.1: “Left ramp” membership function: $\mu(x) = 1$ when $x < 0$, linearly decreasing on $[0, 0.33]$, and $\mu(x) = 0$ when $x > 0.33$.

```

6      activation: General
7      rule: if PhysicalInclination is low and MentalInclination is low
8              then MannersClass is high
9      rule: if PhysicalInclination is low and MentalInclination is low
10         then Hunting is low

```

RuleBlock declaration specifies what is the exact variant of fuzzy controller we are going to use, the Mamdani [9] one or Sugeno [19] one. In the scope of this work we will be using Mamdani controllers exclusively.

5.3 Pagmo library for evolutionary computations

Authors of the Pagmo library designed a very high-level API for the evolutionary computation, which can be completely summarized in the following code snippet:

```

1  // declare the problem to solve
2  pagmo::problem prob(pm_problem{});
3
4  // declare the algorithm to use
5  pagmo::algorithm algo(pagmo::sade(100));
6
7  // declare the population to use
8  pagmo::archipelago archi(16u, algo, prob, 20u);
9
10 // work

```

```

11 archi.evolve(10);
12 archi.wait_check();
13
14 for (const auto& isl : archi)
15 {
16     const auto& champion = isl.get_population().champion_x();
17     // for example, print the champion.
18 }

```

Pagmo library includes a large amount of already implemented algorithms which brings us two benefits:

1. there's no need to implement them from scratch
2. it's easy to experiment as we can change the algorithm just by changing the function name to use in the `pagmo::algorithm` object creation.

Pagmo uses a Generalized Island Model to perform computations in parallel [7], and because of that instead of the “population” the base terminology is “archipelago”. An archipelago consists of “islands”, each containing a separate population to evolve.

After we run an evolution, we can visit every island and check the final result of the population evolution, including getting the best specimen, called “champion” in Pagmo.

5.4 Implementing the method using fuzzylite and pagmo libraries

To implement our solver in this computational framework, first we need to declare our own “problem” class compatible with `pagmo::problem` class requirements.

For this, in the span of this work, it is enough for us to declare the following structure with two member functions:

```

1 struct pm_problem {
2     std::pair<pagmo::vector_double, pagmo::vector_double> get_bounds() const
3     {
4         return { {0., 0.}, {1., 1.} };
5     }
6
7     // Implementation of the objective function.
8     pagmo::vector_double fitness(const pagmo::vector_double& dv) const
9     {
10        // ...implementation of the fitness function...
11    }
12 };

```

The function `get_bounds` describes the range of values from which the specimens would be generated.

In the example above, we declare that every specimen is a vector of two real values between 0 and 1 inclusive. This is the setup for the “Trivial case” experiment from the

section 6. Every element of the vector would be used as an inclination value in the fuzzy inference stage. Depending on the amount of inclination values we require, we specify vectors of according length as return values of this function.

The function **fitness** calculates the fitness value of the given specimen. In a classical evolutionary optimization problems, this fitness function is a pure mathematical function easy to calculate, so the algorithm itself presents the main computational challenge for the computing device. In our problem, however, this function contains the implementation of the control loop described in the section 4.2.

```

1 pagmo::vector_double fitness(const pagmo::vector_double& dv) const
2 {
3     // Inclinations is an alias for std::tuple<double, double>
4     const Inclinations specimen{ dv[0], dv[1] };
5
6     auto engine = init();
7
8     return { simulate(specimen, engine.get()) };
9 }

```

This function does three things:

1. converts the specimen from the raw data provided by Pagmo to the vector of inclinations we use in our fuzzy controller. For simplicity the mapping is direct: values of genes and values of inclinations both belong to $[0, 1] \in \mathbb{R}$.
2. initializes the fuzzy controller ⁵
3. runs the full simulation

Fuzzy controller initialization is a by-the-book copy of the code from fuzzylite documentation.

```

1 std::unique_ptr<fl::Engine> init()
2 {
3     // Initialize the engine
4     std::string path{ "C:\\projects\\pm_solver\\Trivial.fl" };
5     std::unique_ptr<fl::Engine> engine{ fl::FllImporter().FromFile(path) };
6
7     // Checking for errors in the engine loading.
8     std::string status;
9     if (not engine->isReady(&status))
10    {
11        throw fl::Exception("[engine_error] engine is not ready:\n" + status);
12    }
13
14    return engine;
15 }

```

⁵In the final implementation of the solver, initialization of the engine is performed in a more optimized way, only once at the launch of the program. Instead of initializing the fuzzy controller every time for each instance of the problem in a highly parallel Pagmo run, we make an individual copy of the already prepared controller for each specimen to use.

The main point is that the specification of input and output variables and fuzzy rules is kept as a separate text file loaded at runtime.

The simulation is an implementation of the control loop from the section 2

```

1 double simulate(const Inclinations& inclinations , fl::Engine* engine)
2 {
3     // Initialize a character
4     Stats stats{ 0.0, 0.0, 0.0, 0.0 };
5
6     // Set inclinations
7     engine->getInputVariable("PhysicalInclination")
8         ->setValue(std::get<0>(inclinations));
9     engine->getInputVariable("MentalInclination")
10        ->setValue(std::get<1>(inclinations));
11
12     for (int i = 0; i < T; ++i)
13     {
14         single_step(stats , engine);
15     }
16
17     return fitness(stats);
18 }

```

We prepare the character with the starting characteristics, load the fuzzy controller with the values of inclinations and then repeat the action choice and application T times.

```

1 void single_step(Stats& stats , fl::Engine* engine)
2 {
3     // Choose an action based on the current stats and inclinations
4     std::string chosen_action_name = choose_action(engine , stats);
5
6     // Apply the effects of the chosen action
7     Stats stats_diff = actions.at(chosen_action_name);
8     // sum_stats is just a helper for summing two std::tuple values
9     stats = sum_stats(stats , stats_diff);
10 }

```

This function for simplicity of implementation references the statically created global dictionary **actions** which maps action names to the changes in characteristics. With this approach once we know the name of the chosen action we can extract the characteristics changes vector and apply it to the current character state via summation.

Below is an example declaration of **actions** from the section 6. Four real values bound to each action name are changes in the four characteristics of the current character (changes in the current state of the agent).

```

1 const std::unordered_map<
2     std::string , // job name
3     Stats // stat changes after taking this action
4 > actions{
5     { "Hunting" , { 0.00, 0.01, 0.00, -0.01 } },
6     { "Lumberjack" , { 0.02, 0.00, 0.00, -0.02 } },

```

```

7   { "ScienceClass", { 0.0, 0.0, 0.02, 0.0 } },
8   { "MannersClass", { 0.0, 0.0, 0.0, 0.02 } },
9 };

```

An implementation of the **choose_action** is very technical but it boils down to just two things:

1. run the fuzzy controller loaded with the current values of stats and inclinations
2. figure out what output variable got the highest defuzzified value.

```

1  std::string choose_action(fl::Engine* engine, const Stats& stats)
2  {
3      // Load the specimen into the engine
4      // assume that inclinations are already set
5
6      engine->getInputVariable("strength")->setValue(std::get<0>(stats));
7      engine->getInputVariable("constitution")->setValue(std::get<1>(stats));
8      engine->getInputVariable("intelligence")->setValue(std::get<2>(stats));
9      engine->getInputVariable("refinement")->setValue(std::get<3>(stats));
10
11     // Get action priorities
12     engine->process();
13
14     const auto& output_vars = engine->outputVariables();
15     auto it = std::max_element(
16         output_vars.begin(), output_vars.end(),
17         [](const auto* a, const auto* b) {
18             // defaulting to 0 if the value is NaN
19             const auto left_priority = std::isnan(a->getValue())
20                 ? 0.0
21                 : a->getValue();
22             const auto right_priority = std::isnan(b->getValue())
23                 ? 0.0
24                 : b->getValue();
25
26             return left_priority < right_priority;
27         }
28     );
29
30     // Either return the name of the action with the highest priority,
31     // or the first action if no rules fired (i.e., all priorities are 0).
32     std::string chosen_action_name = (it != output_vars.end())
33         ? (*it)->getName()
34         : (*output_vars.begin())->getName();
35
36     return chosen_action_name;
37 }

```

In the listing, we are setting four characteristics — this is the setup for the trivial case from the section 6.

After we finish the simulation, we calculate the fitness. Below is an example from the trivial case, where the goal state is just for one of the characteristics to become higher than the threshold value 0.05.

```

1  /** the lower the better (conforming to pagmo2 conventions) */
2  double fitness(const Stats& stats)
3  {
4      // demo fitness: desirable refinement is 0.05+
5      return 0.05 - std::get<3>(stats);
6  }

```

This completes the full code for the solver.

5.5 Using the implemented solver

Internal beauty of the code and external configurability were not the goals of the implementation, so to prepare the solver for the given problem a set of changes has to be done to the code itself.

According to the definitions in 2, to specify the problem we need to specify the following parameters:

1. a list of n numerical characteristics
2. a list A of m possible actions
3. a list of changes for characteristics for each action
4. a fitness function Φ mapping the characteristics to a single real number, representing the goal state
5. a planning horizon T

All these items are hardcoded in the implementation, directly as types and constants in the C++ source code.

In addition to that, for the method explored in this work to work we need to specify the following:

1. a list of q inclinations
2. fuzzy variables for all the inclinations and characteristics, with all their fuzzy terms
3. fuzzy variables which represent the priorities of each action, with all their fuzzy terms
4. fuzzy rules binding the inclinations, characteristics and action priorities

The shape of the vector of inclinations is being specified as a type in the source code, but all the fuzzy variables and rules are expressed in a DSL of fuzzylite in a separate file.

The configuration for the fuzzy controller, specifically, the membership functions for the fuzzy terms, aggregation, defuzzification, conjunction, disjunction, implication operators, can be seen as hyperparameters for the solver itself detached from the particular problem instance to solve.

Finally, we have an option to choose from the built-in evolutionary algorithms in pagmo library and a configuration of the archipelago of populations, which are also part of the hyperparameters.

Having corrected the code to specify all the above settings, the code for solver just compiles and runs as an executable without any arguments.

Chapter 6

Experiment 1: Trivial case

First case has been used solely for assessing the possibility of constructing the system at all. It is too small to be a useful example of problems solvable by the proposed solver.

We define 4 numerical attributes: **Strength**, **Constitution**, **Intelligence** and **Refinement**.

These 4 attributes are changed by 4 mutually exclusive actions, listed in the table 6.1.

Table 6.1: Trivial case actions

Action	Attribute changes			
	Strength	Constitution	Intelligence	Refinement
Hunting	0.00	0.01	0.00	-0.01
Lumberjack	0.02	0.00	0.00	-0.02
ScienceClass	0.00	0.00	0.02	0.00
MannersClass	0.00	0.00	0.00	0.02

We will run this problem on a goal states reachable in 4 steps, meaning, our T is 4 for a trivial case.

This scenario represents a trivial case, with only 4^3 possible action sequences—a total of 64. The small state space allows for exhaustive enumeration and manual verification of results. This case serves to validate the correctness of the implementation and the fuzzy controller, as the system’s behavior can be easily traced and analyzed by hand.

From the list of attributes and actions we synthesize two inclinations: **Physical Inclination** and **Mental Inclination**, which represent, correspondingly, “an inclination to perform actions related to physical attributes improvement” and “an inclination to perform actions related to mental attributes improvement”.

The full text of the fuzzy controller for the trivial case, which binds two inclinations and four actions, looks as follows:

```
1 Engine: Trivial
2
3 # Inclinations
4
5 InputVariable: PhysicalInclination
6   enabled: true
7   range: 0 1.000
```

```

8   lock-range: false
9   term: tiny Ramp 0.330 0.000
10  term: low Triangle 0.000 0.330 0.670
11  term: high Triangle 0.330 0.670 1.000
12  term: highest Ramp 0.670 1.000
13
14  InputVariable: MentalInclination
15      enabled: true
16      range: 0 1.000
17      lock-range: false
18      term: tiny Ramp 0.330 0.000
19      term: low Triangle 0.000 0.330 0.670
20      term: high Triangle 0.330 0.670 1.000
21      term: highest Ramp 0.670 1.000
22
23  # Action Priorities
24
25  OutputVariable: Hunting
26      enabled: true
27      range: 0.000 1.000
28      lock-range: false
29      aggregation: Maximum
30      defuzzifier: Centroid 100
31      default: nan
32      lock-previous: false
33      term: low Ramp 1.000 0.000
34      term: high Ramp 0.000 1.000
35
36  OutputVariable: Lumberjack
37      enabled: true
38      range: 0.000 1.000
39      lock-range: false
40      aggregation: Maximum
41      defuzzifier: Centroid 100
42      default: nan
43      lock-previous: false
44      term: low Ramp 1.000 0.000
45      term: high Ramp 0.000 1.000
46
47  OutputVariable: ScienceClass
48      enabled: true
49      range: 0.000 1.000
50      lock-range: false
51      aggregation: Maximum
52      defuzzifier: Centroid 100
53      default: nan

```

```

54  lock—previous: false
55  term: low Ramp 1.000 0.000
56  term: high Ramp 0.000 1.000
57
58  OutputVariable: MannersClass
59  enabled: true
60  range: 0.000 1.000
61  lock—range: false
62  aggregation: Maximum
63  defuzzifier: Centroid 100
64  default: nan
65  lock—previous: false
66  term: low Ramp 1.000 0.000
67  term: high Ramp 0.000 1.000
68
69
70  # Rules
71
72  RuleBlock: mamdani
73  enabled: true
74  conjunction: Minimum
75  disjunction: Maximum
76  implication: AlgebraicProduct
77  activation: General
78  rule: if PhysicalInclination is low and MentalInclination is low
79      then MannersClass is high
80  rule: if PhysicalInclination is low and MentalInclination is low
81      then Hunting is low
82  rule: if MentalInclination is high and PhysicalInclination is low
83      then ScienceClass is high
84  rule: if MentalInclination is high and PhysicalInclination is low
85      then Lumberjack is low
86  rule: if PhysicalInclination is high and MentalInclination is low
87      then Lumberjack is high
88  rule: if PhysicalInclination is high and MentalInclination is low
89      then ScienceClass is low
90  rule: if MentalInclination is high and PhysicalInclination is high
91      then Hunting is high
92  rule: if MentalInclination is high and PhysicalInclination is high
93      then MannersClass is low

```

Correspondingly, the goal state is expressed as the following fitness function:

```

1  double fitness(const Stats& stats)
2  {
3      // Stat of index 3 is refinement
4      return 0.07 - std::get<3>(stats);

```

In the 4-tuple **Stats** the element at index 3 (zero-based) is a Refinement attribute, so this fitness function expresses the goal “Refinement must be more than 0.07”.

Given the table of action effects and the goal, it’s obvious that the only correct sequence of actions which can reach this goal is an action “MannersClass” repeated 4 times in a row, as this action is the only way to get an increase in the “Refinement” attribute.

The archipelago has been configured in the following manner:

```
1 pagmo::algorithm algo(pagmo::sade(100));
2 pagmo::archipelago archi(16u, algo, prob, 20u);
```

Four arguments to the `pagmo::archipelago` constructor are number of islands, algorithm to use, problem to solve and size of the population on each island. These are the slice of the hyperparameters which are related to evolutionary optimizations.

The algorithm that has been chosen (`pagmo::sade`) is an instance of Self-adaptive Differential Evolution algorithm, jDE variant [3] [5], for no reason other than being mentioned the first in the Pagmo documentation examples.

Running this system with the above setup results in 16 islands all evolving to the specimen which successfully reach the goal.

The following listing is an example listing of champions across the archipelago — due to stochastic nature of evolutionary optimization, every run will provide different specimen.

```
{0.292813, 0.31037}    fitness -0.03
{0.223563, 0.274446}  fitness -0.03
{0.667839, 0.0552651} fitness -0.03
... 12 more ...
{0.036359, 0.0444171} fitness -0.03
```

Due to the deterministic nature of the problem, we can get the exact list of actions from the specimen by simply calling the `simulate` function (explained in the section 5.4) with the inclination values gotten from the archipelago champions.

In this case, the list of actions is correctly inferred as 4 instances of “MannersClass”.

If we change the goal to check the “Constitution”, then the system converges to an action sequence of 4 instances of a “Hunting” action, with the following specimen examples:

```
{0.379914, 0.929484} fitness 0.01
{0.334496, 0.832509} fitness 0.01
{0.96071, 0.470057}  fitness 0.01
... 12 more ...
{0.516274, 0.708208} fitness 0.01
```


6.1 Discussion of the results

This basic trivial case confirms that we are indeed able to construct the hybrid fuzzy-genetic system which is able to produce optimal sequences of actions which lead to the goal state.

Using the fuzzy controller came out more complicated than it could be seen from the theory alone. While the target of the work was reducing the search state space, the amount of parameters in the fuzzy controller exploded the hyperparameters space instead, as by different configuration of the fuzzy rules and fuzzy variables we can change the behavior of the specimen.

It can be seen that if we exclude the current state of the specimen from the fuzzy rules, we trivialize the trajectories, reducing them to repetition of the same action T times. While this does not simplify the optimization step, as it is assumed that though (3.2), q is still large enough for the brute-force enumeration to be intractable, it leaves us with action sequences intuitively unfit as solutions for any realistic nontrivial goal states.

If we setup the fuzzy action-prioritizing rules in such a way that they would indeed use the current state of the specimen, we do turn the problem into the control one with non-trivial solutions, but at the same time we end up having to specify not only at least one rule per each action priority, but also at least one rule per each *term* per *each input variable*, which starts competing with the complexity of the problem we are trying to solve by this method in the first place.

In the discussion of reducing the dimensionality of the problem, one should not forget the actual issue which explodes the dimensionality of the problem. While on the surface the method explored in this work is based on the inequality (3.2) and the amount of inclinations seems to be the target of discussions, the actual solution to the origin problem is a *list of actions*, and the true reason for dimensionality explosion is the length T of the list of actions to find and the amount of actions to be considered at each step.

It also can be seen that the configuration of the fuzzy controller and a tiebreaking rule is paramount for getting the solution. Intuitively by construction it can be expected that, if we set some “Physical” attribute as a goal, we expect that the specimen with high “Physical Inclination” and low “Mental Inclination” will be the only possible solution, but it’s not the case.

In the example results for the “Strength must be higher than 0.07” case above, we can see a winning specimen with “Physical Inclination” being as low as 0.379914 and “Mental Inclination” being as high as 0.929484, which is a situation completely opposite to the expectations. If we single-step the evolution process for this specimen, we can see that on the step of choosing an action it gets actions with identical priorities, but opposite effects.

For a visual example, this is how the text report looks like if we introduce it in the program appropriately:

```
island champion: {0.379914, 0.929484}
```

```
Step 1:
```

```
Comparing Hunting with value 0.66665 and Lumberjack with value 0.33335
```

```
Comparing Hunting with value 0.66665 and ScienceClass with value 0.66665
```

```
Comparing Hunting with value 0.66665 and MannersClass with value 0.33335
```

```
Chosen action: Hunting
```

Three following steps are not shown because they are identical to this one. We can see that priorities of the “Hunting” and “ScienceClass” actions were inferred as identical, and the “Hunting” action has been chosen solely by the reason of being the first in the list of actions compared to “ScienceClass”.

From one perspective we can see it as a deficiency in an algorithm and implement a more robust tiebreaker, but on the other hand we should not forget that the solution we seek is the list of actions leading to the goal state, not the inclination values which are essentially transient intermediaries representing paths in the actions space.

Chapter 7

Experiment 2: Base control case

24 characteristics, 25 actions with varied amount of steps.

This case is the base case, as it introduces enough complexity to test the proposed approach and at the same time compare it with classical approaches.

A decision tree of the size 25^{100} is already too large to be completely enumerated, and as we show later in this section, we go twelve times deeper than that.

In this case we copy all the numerical attributes of the character sans one from the original *Princess Maker 2* problem, and all the actions with their effects on these attributes.

The attributes are represented as the following tuple:

```
1 using Stats = std::tuple<
2     int, //0 strength
3     int, // constitution
4     int, // intelligence
5     int, // refinement
6     int, //4 charisma
7     int, // morality
8     int, // faith
9     int, // sin
10    int, // sensitivity
11
12    int, //9 combat skill
13    int, // combat attack
14    int, // combat defense
15    int, // magic skill
16    int, // magic attack
17    int, //14 magic defense
18    int, // decorum
19    int, // artistry skill
20    int, // eloquence
21    int, // cooking skill
22    int, //19 cleaning skill
23    int // temperament
24 >;
```

The effects of actions are listed in the table 7.2 for reference.

Applying the methodology explored in this work, we synthesize the 5 inclinations out of all the attributes and actions:

```

1 using Inclinations = std::tuple<
2     double, // fighting
3     double, // magic
4     double, // housekeeping
5     double, // artistry
6     double // sinfulness
7 >;

```

Each inclination represents literally an inclination towards performing a particular set of actions, which we express as a fuzzy rules in a fuzzy controller.

Compared to the trivial case from the chapter 6, now we will use the rules which include both inclination values and the current attribute values of the character. This allows the agent to actually “choose” actions to perform depending on the “actual needs”.

The full code for the fuzzy controller is in the appendix A.

This is a case representing the full complexity we want to explore in the scope of this work, as it was never intended to replicate the original gameplay. The other way around, we wanted to strip away all the unnecessary things from the game to leave only the essential generic problem.

With such an elaborate set of attributes we can formulate complicated fitness rules, for example, the “High General” ending ¹:

1. Intelligence attribute > 500 and > than Sensitivity attribute
2. Morality attribute > 30
3. Faith attribute > 300 and > than Sensitivity attribute
4. “Fighting reputation” which is a sum of all fighting-related stats and skills, > 421

The following piece of code is an example of encoding this ending into a fitness function:

```

1 double fitness(const Stats& stats)
2 {
3     const int fighter_reputation = std::get<0>(stats) + std::get<1>(stats)
4     + std::get<9>(stats) + std::get<10>(stats) + std::get<11>(stats);
5
6     // if sensitivity is higher than intelligence or faith ,
7     // return an absolute instant loss
8     if (std::get<8>(stats) >= std::get<2>(stats)
9         || std::get<8>(stats) >= std::get<6>(stats)) {
10         return std::numeric_limits<double>::max();
11     }
12
13     double fitness_value = 0.0;
14

```

¹[https://princessmaker.fandom.com/wiki/High_General_Ending_\(PM2\)](https://princessmaker.fandom.com/wiki/High_General_Ending_(PM2))

```

15 // add penalty for intelligence below 500
16 if (std::get<2>(stats) < 500) {
17     fitness_value += (500 - std::get<2>(stats));
18 }
19
20 // add penalty for morality below 30
21 if (std::get<5>(stats) < 30) {
22     fitness_value += (30 - std::get<5>(stats));
23 }
24
25 // add penalty for faith below 300
26 if (std::get<6>(stats) < 300) {
27     fitness_value += (300 - std::get<6>(stats));
28 }
29
30 // add penalty for fighter reputation below 421
31 if (fighter_reputation < 421) {
32     fitness_value += (421 - fighter_reputation);
33 }
34
35 return fitness_value;
36 }

```

When we run it on $T = 1200$ which is already far outside of any classical algorithms, we get the following result (an output verbatim from the reference implementation of the solver):

Stats changes:

```

str: +0, con: +0, int: +507,
ref: +0, cha: +0, mor: +102,
fai: +612, sin: -204, sen: -588
cs: +330, ca: +165, cd: +165,
ms: +0, ma: +0, md: +408,
dec: +0, art: +0, elo: +0,
coo: +294, cle: +294, tem: +294

```

Simulation Result:

Fitness: 0

With the runtime of 12 minutes 30 seconds on the standard home PC machine ².

To decode the raw output above, we received a specimen with the:

1. “Intelligence” of 507,
2. “Morality” of 102,
3. “Faith” of 612,

²Intel(R) Core(TM) i7-9700 CPU @ 3.00GHz with 32 GB DDR4 RAM, NVIDIA(R) GeForce(TM) RTX 2060

4. “Sensitivity” of -588 which is guaranteed to be less than both intelligence and faith
5. “Combat skill” 330, “Combat attack” 165 and “Combat defense” 165 summing up to more than 421.

These are exactly the stats we asked for in the goal function, which the fitness score of 0 represents.

And we decode the full plan from the inclination values which is the solution bringing us to these attribute values:

1. TheologyClass x 383
2. FencingClass x 106
3. FightingClass x 106
4. TheologyClass - FencingClass - FightingClass x 24
5. TheologyClass
6. FencingClass - FightingClass x 2
7. TheologyClass - FencingClass - FightingClass x 33
8. Housework x 294
9. Church x 63
10. Church - ScienceClass x 3
11. Church
12. Church - ScienceClass x 6
13. Church
14. Church - ScienceClass x 7
15. Church
16. Church - ScienceClass x 6
17. Church
18. Church - ScienceClass x 7
19. Church
20. Church - ScienceClass x 4
21. Church

The runtime is dependent on the complexity of the fitness function.

If we remove the check for the fighter reputation, lines 28-31 from the listing 7, it cuts the runtime to 6 minutes 30 seconds to reach the fitness 0.

If we remove two checks, lines 23-31 from the same listing, we get the runtime of 5 minutes 50 seconds.

Then, removing additionally the check for morality, lines 18-21, we get the runtime of 6 minutes 10 seconds.

And, finally, if we reduce the fitness function to only the check for sensitivity being lower than intelligence and faith, that is, keep only the lines 6-11 and 33 of the listing 7, the runtime becomes 3 minutes 20 seconds.

We can summarize the dependency of the runtime on the complexity of the fitness function calculation in the table 7.1

Table 7.1: Runtime versus fitness function complexity, baseline case experiment

Requirements clauses amount	Runtime on a reference machine
full “General” ending goal	12 m 30 s
no reputation check	6 m 30 s
no reputation, no faith checks	5 m 50 s
no reputation, no faith, no morality checks	6 m 10 s
no attribute checks, only the predicate	3 m 20 s

7.1 Discussion of the results

First note we must make immediately is that the T is actually dependent on the goal. When we have a real goal we want to evaluate, it has less significance, but when experimenting, it’s possible to set the goal which is completely unreachable by the solver — for example, if we demand a value of an attribute to be 100 after 10 steps ($T = 10$), but every action increasing this attribute increases it by 1.

Pagmo is a highly optimized library so it detects when the chosen evolutionary algorithm converges and stops the evolution returning the fitness reached. This prevents us from naively benchmarking the performance of the solver because when we change the goal, runtime can change arbitrarily. Essentially, adding more conditions — making the fitness calculation depending on more attributes — tend to increase the runtime, as the evolution starts to converge slower.

Then, we received a result for $T = 1200$ which is already outside of all the initial expectations in 12 minutes, a time minuscule compared to modern demands in scientific computations. It is clear that we successfully reduced the combinatorically hard problem to something with at least a polynomial time. More than that, given the nature of a fuzzy controller, on every step it performs the same amount of calculations, which means that each simulation for an individual specimen is $O(T)$, and the computational complexity completely shifts to the evolutionary algorithm, which is reflected in our observations over the dependence of a runtime on a complexity of the goal function.

When writing the inference rules of the fuzzy controller, there is a choice of two ways.

Table 7.2: Action effects per day in Princess Maker 2 (nonzero shown; + indicates increase).

Action	STR	CON	INT	REF	CHA	MOR	FAI	SIN	SEN	CS	CA	CD	MS	MA	MD	DEC	ART	ELO	COO	CLE	TEM
Hunting		+1		-1				+1		+1											
Lumberjack	+2			-2					-2										+1	+1	+1
Housework									+1												
Babysitting					-1																
Church						+1	+2	-2													
Farming	+1	+1		-1																	
Innkeeping										-1										+1	
Restaurant										-1									+1		
Salon	-1								+1												
Masonry		+2			-1																
Graveyard					-1				+1						+1			+1	+1		
Bar			-2																		
Tutoring					-1	+1															
SleazyBar					+2	-3	-3	+2													-1
Cabaret			-1	-2	+3			+1													-1
DanceClass	+1				+1												+1				
FencingClass										+1	+1										
FightingClass										+1		+1									
MagicClass													+1	+2							
PaintingClass									+1								+2				
PoetryClass	+1		+1	+1					+1								+1				
StrategyClass	+2		+2						-1	+1											
ScienceClass	+2		+2				-1								-1						
TheologyClass	+1		+1				+1								+1						
MannersClass				+1												+1					

First way, we can express the domain knowledge by writing rules “from the inclinations to actions”, that is, for each term of every inclination (input variable) describe what impact it should have on the priority of the actions. This option has two drawbacks:

1. it scales with the amount of inclinations times amount of their terms
2. we risk missing out some of the actions if we made a mistake in synthesizing the list of inclinations

This problem becomes much more significant once we start using the attributes in the rules, as the number of attributes is larger than the number of inclinations due to (3.2).

Second way, we can express the domain knowledge by writing rules “from the actions back to inclinations”, that is, for every term of every action (output variable) we describe what combination of inclination and attribute values can lead to it. This option has its drawbacks, too:

1. it scales with the amount of actions times amount of their terms
2. we risk missing out some of the combinations of inclinations and attributes

If we go this way and start adding more rules than the absolute minimum we still risk getting into the combinatorial explosion.

Still, the drawbacks listed here do not mean that there are absolute requirements on the amount of rules we have to write. The more specific we would be in writing the rules, the more correct the action selection is expected to be.

One of the observed situation was skipping the input variables because of over-specification of terms.

Let’s say we specify our “artistry” and “housekeeping” inclinations as follows, with four terms *tiny*, *low*, *high*, *highest*:

```

1 InputVariable: InclinationArtistry
2   enabled: true
3   range: 0 1.000
4   lock-range: false
5   term: tiny Ramp 0.330 0.000
6   term: low Triangle 0.000 0.330 0.670
7   term: high Triangle 0.330 0.670 1.000
8   term: highest Ramp 0.670 1.000
9
10  # (identical definition for InclinationHousekeeping)

```

And we also have a block of rules like following:

```

1 rule: if InclinationArtistry is high then MannersClass is high
2 rule: if InclinationHousekeeping is high then MannersClass is low

```

If we have a specimen with **InclinationArtistry** = 0.9 and **InclinationHousekeeping** = 0.6, then, unexpectedly to us, priority of **MannersClass** would be low, not high. Because the membership function for the term “high” will emit higher value for “housekeeping” inclination than for the “artistry”. As we don’t have a rule for **InclinationArtistry**

is **highest** the values *higher* than specified in the rule, despite being “obviously” fitting the rule, are actually ignored instead.

To resolve this problem it’s more beneficial to actually reduce the amount of terms, removing the ones not used in the rules.

Chapter 8

Conclusions and Future Work

To sum it up, the main result of this work is that the hybrid fuzzy-genetic solver is able to function as intended (instead of: “actually works”). It produces action sequences with fitness 0, meaning they fully satisfy the goal predicate. That was the critical test of whether the idea is viable at all, and we passed it.

Performance-wise, the solver depends linearly on the planning horizon T . For example, running $n = 1200$ $T=1200$ steps on a standard home PC took 12 minutes. For modern computing this is effectively instantaneous, and it completely sidesteps the combinatorial explosion of classical approaches. The runtime instead depends mostly on how complex the fitness function is and, to a lesser extent, how the fuzzy controller is written.

Along the way, we found several structural properties:

- The fuzzy controller, while useful for encoding domain knowledge, is a complicated object on its own. Configuring it is already a nontrivial problem.
- The convergence speed of the genetic algorithm is governed by the cost of evaluating the fitness function, not by the size of the action space.

All objectives defined in Section 3.3 have been met: we formalized the problem, implemented a solver, tested it on different scales, and analyzed the results.

A useful comparison is with reinforcement learning. In our case, if we try to set up a tabular RL agent with 23 attributes ranging up to 500 and 27 actions, the Q-table would be 27×500^{23} , which is clearly infeasible (instead of: “obviously impossible”) to store or explore. Even with function approximation, RL still faces the horizon problem: with $\gamma = 0.99$ and $T = 1200$, the learning signal from the final reward is around 6×10^{-6} , practically zero. Raising γ makes the effective horizon too long and increases sample complexity. Our method avoids this by using the known deterministic model to directly simulate trajectories, and by searching in the much smaller space of inclinations $q \in [0, 1]^n$. The trade-off is that we have to supply a fuzzy rule base by hand. In short (instead of: “In other words”): we pay in design effort, but in return we get a tractable solver that actually reaches the goal.

There are limitations. The model used here is a simplification of the full Princess Maker 2 game, and the solver is sensitive to how the fuzzy rules are written. The fuzzy controller should not be seen as a simplification tool but rather as a way to make unsolvable problems solvable at all.

As for future work, there are several directions that go beyond this thesis:

1. Study the behavior of the solver when the number of attributes n and actions

m increase.

2. Extend the solver to non-deterministic environments. The fuzzy part will still work, but decoding the final plan from inclinations will be harder.
3. Benchmark against existing planners and RL-EA hybrids to see where this method fits in the larger ecosystem.
4. Automate the fuzzy rule generation with clustering or learning, to reduce the manual workload.
5. Try other hybridizations, such as RL-assisted evolutionary search or temporal-logic-based planners with domain knowledge.

The final takeaway is that the fuzzy-genetic solver is a proof of concept. It shows that dimensionality reduction and symbolic reasoning can turn an intractable problem into something we can actually solve. The Princess Maker 2 example was just a testbed. The method itself applies to any large-horizon deterministic planning problem, and pushing this further would be a natural path for doctoral-level research.

To sum it up, we achieved the following results:

1. a reference solver implementing the method described in this work is actually able to produce the solutions with fitness 0, meaning, the ones which fully satisfy the predefined goal predicate;
2. performance-wise, the solver depends on T linearly, effectively completely eliminating the computational complexity dependent on the amount of actions to execute;
3. discoveries has been done related to the actual complexity of encoding the domain knowledge as a fuzzy controller;
4. it has been determined that the speed of convergence is actually dependent on the complexity of the fitness function calculation;

The codebase of the reference solver is located at the GitHub repository in the Internet accessible by the URL <https://github.com/hijarian/urv-mesiia-2025-tfm-code>.

With this, our objectives stated in the section 3.3 has been reached and we can give a brief summary on the further research spanning outside the scope of this particular work.

A game so mechanically interdependent as *Princess Maker 2* is too complex to solve using such a straightforward model. A lot of simplifications were needed to arrive at a coherent model of the deterministic planning problem with a stateful agent. Nevertheless, the hybrid solver built in the scope of this work have shown that it successfully determines a solution given that all the domain knowledge is accurately encoded in the fuzzy controller, and as such, it can be used to solve any high-dimensional deterministic planning problem which can be modeled according to our definitions from the section 2. This constitutes the core value of this work.

We have shown that the hybrid fuzzy-genetic system described in this work has a linear complexity in relation to the amount of planning steps T assuming that executing T steps is a requirement. Computing $T = 1200$ steps in 12 minutes is effectively instantaneous nowadays so there's a little benefit in a further research of the computational complexity *in relation to* T . However, a research on how the computational complexity of this solver depends on amount of attributes n (2.1) and amount of actions m (2.2) is still needed to be done.

In the span of this work, only two cases distinct in the number of actions m were explored, and the more thorough exploration of the parameter space is left for a dissertation-level research.

An additional path of exploration is researching how well this method can be applied to problems where the effects of actions are non-deterministic. While lifting the restriction of determinism does not affect the fuzzy inference at all, it presents challenges in decoding the solution sequence of actions from the vector of inclination values.

Now, first, let us make a conclusion tangential to the origin problem but related to the chosen theoretical toolset. While the fuzzy controller can indeed be seen as a tool to formalize the decision making using the expert knowledge in the given domain, it is too complicated mechanism *by itself* to help reducing the complexity of the problem it's solving. One should treat it not as a tool which helps make complicated problems simpler but which, hopefully, makes unsolvable ones solvable at all, as proper configuration of the fuzzy controller is already a problem in itself.

Second, the proof-of-concept solver implemented and discussed in this work is obviously an instance of an automated planner, for which an extensive framework of both machinery and benchmarking exists. The logical next step would be to determine a correct place of this method in the existing body of knowledge about the automated planners and perform formal comparisons of performance with some other state-of-the-art methods.

Another work which exploits the idea of using the domain knowledge in the solution process exists, but it uses the Linear Temporal Logic instead, and utilizes the existing PDDL automatic planning machinery [8].

The paper [16] briefly mentioned in the section 4.3, is an possible alternative hybrid system, where the target of evolution is a reinforcement learning process.

Both these are examples of possible further research in the area of hybrid systems solving the high-dimensional planning problems of large horizon using the domain knowledge.

Bibliography

- [1] Hans-Georg Beyer and Hans-Paul Schwefel. “Evolution Strategies – A Comprehensive Introduction”. In: *Natural Computing* 1.1 (2002), pp. 3–52. DOI: [10.1023/A:1015059928466](https://doi.org/10.1023/A:1015059928466).
- [2] Francesco Biscani and Dario Izzo. “A parallel global multiobjective framework for optimization: pagmo”. In: *Journal of Open Source Software* 5.53 (2020), p. 2338. DOI: [10.21105/joss.02338](https://doi.org/10.21105/joss.02338). URL: <https://doi.org/10.21105/joss.02338>.
- [3] Janez Brest et al. “Self-adapting Control Parameters in Differential Evolution: A Comparative Study on Numerical Benchmark Problems”. In: *IEEE Transactions on Evolutionary Computation* 10.6 (Dec. 2006), pp. 646–657. DOI: [10.1109/TEVC.2006.872133](https://doi.org/10.1109/TEVC.2006.872133). URL: <https://doi.org/10.1109/TEVC.2006.872133>.
- [4] Shi Cheng, Hui Lu, and Xiujuan Lei. “Automated Planning and Scheduling with Swarm Intelligence”. In: *ADVANCES IN SWARM INTELLIGENCE, PT II, ICSI 2024*. Ed. by Y Tan and Y Shi. Vol. 14789. Lecture Notes in Computer Science. 15th International Conference on Advances in Swarm Intelligence (ICSI), Xining, PEOPLES R CHINA, AUG 23-26, 2024. Int Assoc Swarm & Evolut Intellig. 2024, pp. 26–35. ISBN: 978-981-97-7183-7; 978-981-97-7184-4. DOI: [10.1007/978-981-97-7184-4_3](https://doi.org/10.1007/978-981-97-7184-4_3).
- [5] Saber M. Elsayed, Ruhul A. Sarker, and Daryl L. Essam. “Differential Evolution with Multiple Strategies for Solving CEC2011 Real-World Numerical Optimization Problems”. In: *Proceedings of the 2011 IEEE Congress on Evolutionary Computation (CEC)*. New Orleans, LA, USA: IEEE, June 2011, pp. 1041–1048. DOI: [10.1109/CEC.2011.5949732](https://doi.org/10.1109/CEC.2011.5949732). URL: <https://doi.org/10.1109/CEC.2011.5949732>.
- [6] Richard E. Fikes and Nils J. Nilsson. “STRIPS: A New Approach to the Application of Theorem Proving to Problem Solving”. In: *Artificial Intelligence* 2.3–4 (1971), pp. 189–208. DOI: [10.1016/0004-3702\(71\)90010-5](https://doi.org/10.1016/0004-3702(71)90010-5).
- [7] Dario Izzo, Marek Ruciński, and Francesco Biscani. “The Generalized Island Model”. In: *Parallel Architectures and Bioinspired Algorithms*. Ed. by Francisco Fernández de Vega, José Ignacio Hidalgo Pérez, and Juan Lanchares. Berlin, Heidelberg: Springer Berlin Heidelberg, 2012, pp. 151–169. ISBN: 978-3-642-28789-3. DOI: [10.1007/978-3-642-28789-3_7](https://doi.org/10.1007/978-3-642-28789-3_7). URL: https://doi.org/10.1007/978-3-642-28789-3_7.
- [8] Xu Lu et al. “On the exploitation of control knowledge for enhancing automated planning”. In: *Information Sciences* 693 (2025), p. 121666. ISSN: 0020-0255. DOI: <https://doi.org/10.1016/j.ins.2024.121666>. URL: <https://www.sciencedirect.com/science/article/pii/S0020025524015809>.

- [9] Ebrahim H. Mamdani. “Application of Fuzzy Algorithms for Control of Simple Dynamic Plant”. In: *Proceedings of the Institution of Electrical Engineers* 121.12 (1974), pp. 1585–1588. DOI: [10.1049/piee.1974.0328](https://doi.org/10.1049/piee.1974.0328).
- [10] Melanie Mitchell. *An Introduction to Genetic Algorithms*. Cambridge, MA: MIT Press, 1999. ISBN: 9780262631853.
- [11] Juan Rada-Vilela. *The FuzzyLite Libraries for Fuzzy Logic Control*. 2018. URL: <https://fuzzylite.com>.
- [12] S. Kumar Ray. *Soft Computing and Its Applications, Volume II*. Boca Raton, FL: CRC Press, 2014. ISBN: 978-1-4822-5793-9.
- [13] Enrico Scala et al. “Interval-Based Relaxation for General Numeric Planning”. In: *European Conference on Artificial Intelligence*. 2016. URL: <https://api.semanticscholar.org/CorpusID:27984436>.
- [14] Enrico Scala et al. “Subgoalting Techniques for Satisficing and Optimal Numeric Planning”. In: *Journal of Artificial Intelligence Research* 68 (Aug. 10, 2020), pp. 691–752. DOI: [10.1613/JAIR.1.11875](https://doi.org/10.1613/JAIR.1.11875). URL: <https://www.jair.org/index.php/jair/article/view/11875>.
- [15] Hans-Paul Schwefel. *Numerical Optimization of Computer Models*. Chichester, UK: John Wiley & Sons, 1981. ISBN: 9780471099888.
- [16] Yanjie Song et al. “Reinforcement Learning-assisted Evolutionary Algorithm: A Survey”. In: *arXiv preprint arXiv:2308.13420* (2023). Comprehensive taxonomy of RL–EA hybrids.
- [17] Bjarne Stroustrup. *Programming: Principles and Practice Using C++, 3rd Edition*. 3rd ed. Available online. Addison-Wesley Professional, 2024. ISBN: 978-0-13-830868-1. URL: <https://www.informit.com/store/programming-principles-and-practice-using-c-plus-plus-9780138308681>.
- [18] Richard S. Sutton and Andrew G. Barto. *Reinforcement Learning: An Introduction*. 2nd ed. Available online. MIT Press, 2018. ISBN: 978-0-262-03924-6. URL: <http://incompleteideas.net/book/the-book-2nd.html>.
- [19] T. Takagi and M. Sugeno. “Fuzzy Identification of Systems and Its Applications to Modeling and Control”. In: *IEEE Transactions on Systems, Man, and Cybernetics* SMC-15.1 (1985), pp. 116–132. DOI: [10.1109/TSMC.1985.6313399](https://doi.org/10.1109/TSMC.1985.6313399).
- [20] Lotfi A. Zadeh. “Fuzzy Sets”. In: *Information and Control* 8.3 (1965), pp. 338–353. DOI: [10.1016/S0019-9958\(65\)90241-X](https://doi.org/10.1016/S0019-9958(65)90241-X).

Appendix A

Baseline case fuzzy controller

```
1 Engine: Baseline
2
3 # Inclinations
4
5 InputVariable: InclinationFighting
6   enabled: true
7   range: 0 1.000
8   lock-range: false
9   term: low Ramp 1.0 0.0
10  term: high Ramp 0.0 1.0
11 InputVariable: InclinationMagic
12   enabled: true
13   range: 0 1.000
14   lock-range: false
15   term: low Ramp 1.0 0.0
16   term: high Ramp 0.0 1.0
17 InputVariable: InclinationHousekeeping
18   enabled: true
19   range: 0 1.000
20   lock-range: false
21   term: low Ramp 1.0 0.0
22   term: high Ramp 0.0 1.0
23 InputVariable: InclinationArtistry
24   enabled: true
25   range: 0 1.000
26   lock-range: false
27   term: low Ramp 1.0 0.0
28   term: high Ramp 0.0 1.0
29 InputVariable: InclinationSinfulness
30   enabled: true
31   range: 0 1.000
32   lock-range: false
33   term: low Ramp 1.0 0.0
```



```

34     term: high Ramp 0.0 1.0
35
36 # Attributes
37
38 InputVariable: strength
39     enabled: true
40     range: 0 810
41     lock-range: false
42     term: low Ramp 810 0
43     term: high Ramp 0 810
44 InputVariable: constitution
45     enabled: true
46     range: 0 810
47     lock-range: false
48     term: low Ramp 810 0
49     term: high Ramp 0 810
50 InputVariable: intelligence
51     enabled: true
52     range: 0 810
53     lock-range: false
54     term: low Ramp 810 0
55     term: high Ramp 0 810
56 InputVariable: refinement
57     enabled: true
58     range: 0 810
59     lock-range: false
60     term: low Ramp 810 0
61     term: high Ramp 0 810
62 InputVariable: charisma
63     enabled: true
64     range: 0 810
65     lock-range: false
66     term: low Ramp 810 0
67     term: high Ramp 0 810
68 InputVariable: morality
69     enabled: true
70     range: 0 810
71     lock-range: false
72     term: low Ramp 810 0
73     term: high Ramp 0 810
74 InputVariable: faith
75     enabled: true
76     range: 0 810
77     lock-range: false
78     term: low Ramp 810 0
79     term: high Ramp 0 810

```

```

80 InputVariable: sinfulness
81     enabled: true
82     range: 0 810
83     lock-range: false
84     term: low Ramp 810 0
85     term: high Ramp 0 810
86 InputVariable: sensitivity
87     enabled: true
88     range: 0 810
89     lock-range: false
90     term: low Ramp 810 0
91     term: high Ramp 0 810
92
93 # Skills (same as attributes but range is 0-540 instead of 0-810)
94 InputVariable: CombatSkill
95     enabled: true
96     range: 0 540
97     lock-range: false
98     term: low Ramp 540 0
99     term: high Ramp 0 540
100 InputVariable: CombatAttack
101     enabled: true
102     range: 0 540
103     lock-range: false
104     term: low Ramp 540 0
105     term: high Ramp 0 540
106 InputVariable: CombatDefense
107     enabled: true
108     range: 0 540
109     lock-range: false
110     term: low Ramp 540 0
111     term: high Ramp 0 540
112 InputVariable: MagicSkill
113     enabled: true
114     range: 0 540
115     lock-range: false
116     term: low Ramp 540 0
117     term: high Ramp 0 540
118 InputVariable: MagicAttack
119     enabled: true
120     range: 0 540
121     lock-range: false
122     term: low Ramp 540 0
123     term: high Ramp 0 540
124 InputVariable: MagicDefense
125     enabled: true

```

```

126     range: 0 540
127     lock-range: false
128     term: low Ramp 540 0
129     term: high Ramp 0 540
130 InputVariable: Decorum
131     enabled: true
132     range: 0 540
133     lock-range: false
134     term: low Ramp 540 0
135     term: high Ramp 0 540
136 InputVariable: Artistry
137     enabled: true
138     range: 0 540
139     lock-range: false
140     term: low Ramp 540 0
141     term: high Ramp 0 540
142 InputVariable: Eloquence
143     enabled: true
144     range: 0 540
145     lock-range: false
146     term: low Ramp 540 0
147     term: high Ramp 0 540
148 InputVariable: CookingSkill
149     enabled: true
150     range: 0 540
151     lock-range: false
152     term: low Ramp 540 0
153     term: high Ramp 0 540
154 InputVariable: CleaningSkill
155     enabled: true
156     range: 0 540
157     lock-range: false
158     term: low Ramp 540 0
159     term: high Ramp 0 540
160 InputVariable: Temperament
161     enabled: true
162     range: 0 540
163     lock-range: false
164     term: low Ramp 540 0
165     term: high Ramp 0 540
166
167 # Action Priorities
168
169 ## Jobs
170
171 OutputVariable: Hunting

```

```

172     enabled: true
173     range: 0.000 1.000
174     lock-range: false
175     aggregation: Maximum
176     defuzzifier: Centroid 100
177     default: nan
178     lock-previous: false
179     term: low Ramp 0.500 0.000
180     term: medium Triangle 0.000 0.500 1.000
181     term: high Ramp 0.500 1.000
182
183 OutputVariable: Lumberjack
184     enabled: true
185     range: 0.000 1.000
186     lock-range: false
187     aggregation: Maximum
188     defuzzifier: Centroid 100
189     default: nan
190     lock-previous: false
191     term: low Ramp 0.500 0.000
192     term: medium Triangle 0.000 0.500 1.000
193     term: high Ramp 0.500 1.000
194
195 OutputVariable: Housework
196     enabled: true
197     range: 0.000 1.000
198     lock-range: false
199     aggregation: Maximum
200     defuzzifier: Centroid 100
201     default: nan
202     lock-previous: false
203     term: low Ramp 0.500 0.000
204     term: medium Triangle 0.000 0.500 1.000
205     term: high Ramp 0.500 1.000
206
207 OutputVariable: Babysitting
208     enabled: true
209     range: 0.000 1.000
210     lock-range: false
211     aggregation: Maximum
212     defuzzifier: Centroid 100
213     default: nan
214     lock-previous: false
215     term: low Ramp 0.500 0.000
216     term: medium Triangle 0.000 0.500 1.000
217     term: high Ramp 0.500 1.000

```

```

218
219 OutputVariable: Church
220     enabled: true
221     range: 0.000 1.000
222     lock-range: false
223     aggregation: Maximum
224     defuzzifier: Centroid 100
225     default: nan
226     lock-previous: false
227     term: low Ramp 0.500 0.000
228     term: medium Triangle 0.000 0.500 1.000
229     term: high Ramp 0.500 1.000
230
231 OutputVariable: Farming
232     enabled: true
233     range: 0.000 1.000
234     lock-range: false
235     aggregation: Maximum
236     defuzzifier: Centroid 100
237     default: nan
238     lock-previous: false
239     term: low Ramp 0.500 0.000
240     term: medium Triangle 0.000 0.500 1.000
241     term: high Ramp 0.500 1.000
242
243 OutputVariable: Innkeeping
244     enabled: true
245     range: 0.000 1.000
246     lock-range: false
247     aggregation: Maximum
248     defuzzifier: Centroid 100
249     default: nan
250     lock-previous: false
251     term: low Ramp 0.500 0.000
252     term: medium Triangle 0.000 0.500 1.000
253     term: high Ramp 0.500 1.000
254
255 OutputVariable: Restaurant
256     enabled: true
257     range: 0.000 1.000
258     lock-range: false
259     aggregation: Maximum
260     defuzzifier: Centroid 100
261     default: nan
262     lock-previous: false
263     term: low Ramp 0.500 0.000

```

```

264     term: medium Triangle 0.000 0.500 1.000
265     term: high Ramp 0.500 1.000
266
267 OutputVariable: Salon
268     enabled: true
269     range: 0.000 1.000
270     lock-range: false
271     aggregation: Maximum
272     defuzzifier: Centroid 100
273     default: nan
274     lock-previous: false
275     term: low Ramp 0.500 0.000
276     term: medium Triangle 0.000 0.500 1.000
277     term: high Ramp 0.500 1.000
278
279 OutputVariable: Masonry
280     enabled: true
281     range: 0.000 1.000
282     lock-range: false
283     aggregation: Maximum
284     defuzzifier: Centroid 100
285     default: nan
286     lock-previous: false
287     term: low Ramp 0.500 0.000
288     term: medium Triangle 0.000 0.500 1.000
289     term: high Ramp 0.500 1.000
290
291 OutputVariable: Graveyard
292     enabled: true
293     range: 0.000 1.000
294     lock-range: false
295     aggregation: Maximum
296     defuzzifier: Centroid 100
297     default: nan
298     lock-previous: false
299     term: low Ramp 0.500 0.000
300     term: medium Triangle 0.000 0.500 1.000
301     term: high Ramp 0.500 1.000
302
303 OutputVariable: Bar
304     enabled: true
305     range: 0.000 1.000
306     lock-range: false
307     aggregation: Maximum
308     defuzzifier: Centroid 100
309     default: nan

```

```

310     lock—previous: false
311     term: low Ramp 0.500 0.000
312     term: medium Triangle 0.000 0.500 1.000
313     term: high Ramp 0.500 1.000
314
315 OutputVariable: Tutoring
316     enabled: true
317     range: 0.000 1.000
318     lock—range: false
319     aggregation: Maximum
320     defuzzifier: Centroid 100
321     default: nan
322     lock—previous: false
323     term: low Ramp 0.500 0.000
324     term: medium Triangle 0.000 0.500 1.000
325     term: high Ramp 0.500 1.000
326
327 OutputVariable: SleazyBar
328     enabled: true
329     range: 0.000 1.000
330     lock—range: false
331     aggregation: Maximum
332     defuzzifier: Centroid 100
333     default: nan
334     lock—previous: false
335     term: low Ramp 0.500 0.000
336     term: medium Triangle 0.000 0.500 1.000
337     term: high Ramp 0.500 1.000
338
339 OutputVariable: Cabaret
340     enabled: true
341     range: 0.000 1.000
342     lock—range: false
343     aggregation: Maximum
344     defuzzifier: Centroid 100
345     default: nan
346     lock—previous: false
347     term: low Ramp 0.500 0.000
348     term: medium Triangle 0.000 0.500 1.000
349     term: high Ramp 0.500 1.000
350
351 ## Classes
352
353 OutputVariable: DanceClass
354     enabled: true
355     range: 0.000 1.000

```

```

356     lock-range: false
357     aggregation: Maximum
358     defuzzifier: Centroid 100
359     default: nan
360     lock-previous: false
361     term: low Ramp 0.500 0.000
362     term: medium Triangle 0.000 0.500 1.000
363     term: high Ramp 0.500 1.000
364
365 OutputVariable: FencingClass
366     enabled: true
367     range: 0.000 1.000
368     lock-range: false
369     aggregation: Maximum
370     defuzzifier: Centroid 100
371     default: nan
372     lock-previous: false
373     term: low Ramp 0.500 0.000
374     term: medium Triangle 0.000 0.500 1.000
375     term: high Ramp 0.500 1.000
376
377 OutputVariable: FightingClass
378     enabled: true
379     range: 0.000 1.000
380     lock-range: false
381     aggregation: Maximum
382     defuzzifier: Centroid 100
383     default: nan
384     lock-previous: false
385     term: low Ramp 0.500 0.000
386     term: medium Triangle 0.000 0.500 1.000
387     term: high Ramp 0.500 1.000
388
389 OutputVariable: MagicClass
390     enabled: true
391     range: 0.000 1.000
392     lock-range: false
393     aggregation: Maximum
394     defuzzifier: Centroid 100
395     default: nan
396     lock-previous: false
397     term: low Ramp 0.500 0.000
398     term: medium Triangle 0.000 0.500 1.000
399     term: high Ramp 0.500 1.000
400
401 OutputVariable: PaintingClass

```



```

402     enabled: true
403     range: 0.000 1.000
404     lock-range: false
405     aggregation: Maximum
406     defuzzifier: Centroid 100
407     default: nan
408     lock-previous: false
409     term: low Ramp 0.500 0.000
410     term: medium Triangle 0.000 0.500 1.000
411     term: high Ramp 0.500 1.000
412
413 OutputVariable: PoetryClass
414     enabled: true
415     range: 0.000 1.000
416     lock-range: false
417     aggregation: Maximum
418     defuzzifier: Centroid 100
419     default: nan
420     lock-previous: false
421     term: low Ramp 0.500 0.000
422     term: medium Triangle 0.000 0.500 1.000
423     term: high Ramp 0.500 1.000
424
425 OutputVariable: StrategyClass
426     enabled: true
427     range: 0.000 1.000
428     lock-range: false
429     aggregation: Maximum
430     defuzzifier: Centroid 100
431     default: nan
432     lock-previous: false
433     term: low Ramp 0.500 0.000
434     term: medium Triangle 0.000 0.500 1.000
435     term: high Ramp 0.500 1.000
436
437 OutputVariable: ScienceClass
438     enabled: true
439     range: 0.000 1.000
440     lock-range: false
441     aggregation: Maximum
442     defuzzifier: Centroid 100
443     default: nan
444     lock-previous: false
445     term: low Ramp 0.500 0.000
446     term: medium Triangle 0.000 0.500 1.000
447     term: high Ramp 0.500 1.000

```

448

```
449 OutputVariable: TheologyClass
450     enabled: true
451     range: 0.000 1.000
452     lock-range: false
453     aggregation: Maximum
454     defuzzifier: Centroid 100
455     default: nan
456     lock-previous: false
457     term: low Ramp 0.500 0.000
458     term: medium Triangle 0.000 0.500 1.000
459     term: high Ramp 0.500 1.000
```

460

```
461 OutputVariable: MannersClass
462     enabled: true
463     range: 0.000 1.000
464     lock-range: false
465     aggregation: Maximum
466     defuzzifier: Centroid 100
467     default: nan
468     lock-previous: false
469     term: low Ramp 0.500 0.000
470     term: medium Triangle 0.000 0.500 1.000
471     term: high Ramp 0.500 1.000
```

472

473

```
474 # Rules
```

475

```
476 RuleBlock: mamdani
477     enabled: true
478     conjunction: Minimum
479     disjunction: Maximum
480     implication: AlgebraicProduct
481     activation: General
```

482

```
483 # at least one rule for each action priority
484 # We will follow a pattern where high inclination + low attribute/skill
485 # and opposite inclination high and the penalized attribute/skill low -
486 # Jobs
487 rule: if InclinationFighting is high and strength is low then Lumberjack
488 rule: if InclinationArtistry is high and refinement is low
then Lumberjack is low
489 #
490 rule: if InclinationFighting is high and constitution is low then Hunting
491 rule: if InclinationArtistry is high and refinement is low
then Hunting is low
```

```

492  #
493  rule: if InclinationHousekeeping is high and CookingSkill is low then H
494  rule: if InclinationHousekeeping is high and CleaningSkill is low then I
495  rule: if InclinationHousekeeping is high and Temperament is low then Ho
496  rule: if InclinationSinfulness is high and sensitivity is low then House
497  #
498  rule: if InclinationHousekeeping is high and sensitivity is low then Ba
499  rule: if InclinationArtistry is high and charisma is low then Babysittin
500  #
501  rule: if InclinationHousekeeping is high and morality is low then Church
502  rule: if InclinationHousekeeping is high and faith is low then Church
503  rule: if InclinationSinfulness is high and sinfulness is low then Church
504  #
505  rule: if InclinationFighting is high and strength is low then Farmin
506  rule: if InclinationFighting is high and constitution is low then Farmin
507  rule: if InclinationArtistry is high and refinement is low then Farmin
508  #
509  rule: if InclinationHousekeeping is high and CleaningSkill is low
then Innkeeping is high
510  rule: if InclinationFighting is high and CombatSkill is low then Innl
511  #
512  rule: if InclinationHousekeeping is high and CookingSkill is low
then Restaurant is high
513  rule: if InclinationFighting is high and CombatSkill is low then Res
514  #
515  rule: if InclinationArtistry is high and sensitivity is low
then Salon is low
516  rule: if InclinationFighting is high and strength is low then Salon i
517  #
518  rule: if InclinationFighting is high and constitution is low then Mas
519  rule: if InclinationArtistry is high and charisma is low then Masonr
520  #
521  rule: if InclinationMagic is high and MagicDefense is low then Graveya
522  rule: if InclinationArtistry is high and sensitivity is low
then Graveyard is high
523  rule: if InclinationArtistry is high and charisma is low then Graveyar
524  #
525  rule: if InclinationArtistry is high and Eloquence is low
then Bar is high
526  rule: if InclinationHousekeeping is high and CookingSkill is low then
527  rule: if InclinationMagic is high and intelligence is low then Bar is
528  #
529  rule: if InclinationHousekeeping is high and morality is low
then Tutoring is high
530  rule: if InclinationArtistry is high and charisma is low then Tutorin
531  #

```

```

532  rule: if InclinationSinfulness is high and sinfulness is low then SleazyBar
533  rule: if InclinationArtistry is high and charisma is low then SleazyBar
534  rule: if InclinationHousekeeping is high and morality is low
then SleazyBar is low
535  rule: if InclinationHousekeeping is high and faith is low then SleazyBar
536  rule: if InclinationHousekeeping is high and Temperament is low
then SleazyBar is low
537  #
538  rule: if InclinationSinfulness is high and sinfulness is low
then Cabaret is high
539  rule: if InclinationArtistry is high and charisma is low then Cabaret
540  rule: if InclinationMagic is high and intelligence is low then Cabaret
541  rule: if InclinationHousekeeping is high and Temperament is low
then Cabaret is low
542  rule: if InclinationArtistry is high and refinement is low
then Cabaret is low
543  # Classes
544  rule: if InclinationFighting is high and constitution is low
then DanceClass is high
545  rule: if InclinationArtistry is high and charisma is low then DanceClass
546  rule: if InclinationArtistry is high and Artistry is low then DanceClass
547  rule: if InclinationMagic is high then DanceClass is low
548  #
549  rule: if InclinationFighting is high and CombatSkill is low
then FencingClass is high
550  rule: if InclinationFighting is high and CombatAttack is low
then FencingClass is high
551  rule: if InclinationHousekeeping is high then FencingClass is low
552  #
553  rule: if InclinationFighting is high and CombatSkill is low
then FightingClass is high
554  rule: if InclinationFighting is high and CombatDefense is low
then FightingClass is high
555  rule: if InclinationHousekeeping is high then FightingClass is low
556  #
557  rule: if InclinationMagic is high and MagicSkill is low then MagicClass
558  rule: if InclinationMagic is high and MagicAttack is low then MagicClass
559  rule: if InclinationFighting is high then MagicClass is low
560  #
561  rule: if InclinationArtistry is high and sensitivity is low then PaintingClass
562  rule: if InclinationArtistry is high and refinement is low then PaintingClass
563  rule: if InclinationMagic is high and intelligence is low then PaintingClass
564  rule: if InclinationArtistry is high and Artistry is low then PaintingClass
565  rule: if InclinationFighting is high then PaintingClass is low
566  #

```

```

567     rule: if InclinationArtistry is high and refinement is low
then PoetryClass is high
568     rule: if InclinationArtistry is high and sensitivity is low
then PoetryClass is high
569     rule: if InclinationArtistry is high and Artistry is low then PoetryC
570     rule: if InclinationMagic is high and intelligence is low then PoetryC
571     rule: if InclinationFighting is high then PoetryClass is low
572     #
573     rule: if InclinationMagic is high and intelligence is low then Strateg
574     rule: if InclinationFighting is high and CombatSkill is low
then StrategyClass is high
575     rule: if InclinationArtistry is high and sensitivity is low then Strat
576     #
577     rule: if InclinationMagic is high and intelligence is low
then TheologyClass is high
578     rule: if InclinationMagic is high and MagicDefense is low
then TheologyClass is high
579     rule: if InclinationHousekeeping is high and faith is low then Theolog
580     rule: if InclinationSinfulness is high then TheologyClass is low
581     #
582     rule: if InclinationArtistry is high and Decorum is low then Manner
583     rule: if InclinationArtistry is high and refinement is low then Manne
584     rule: if InclinationHousekeeping is high then MannersClass is low
585     #
586     rule: if InclinationMagic is high and intelligence is low then ScienceC
587     rule: if InclinationMagic is high and MagicDefense is low then ScienceC
588     rule: if InclinationHousekeeping is high and faith is low then Scienc

```
